School of Computer Science and Engineering

(Computer Science & Engineering – Artificial Intelligence & Machine Learning)

Faculty of Engineering & Technology

Jain Global Campus, Kanakapura Taluk- 562112

Ramanagara District, Karnataka, India

A Project Report on

Fair Ranking Model for Skill-Based Job Matching

Submitted in partial fulfilment for the award of the degree of

BACHELOR OF TECHNOLOGY IN

COMPUTER SCIENCE AND ENGINEERING

Artificial Intelligence & Machine Learning

Submitted by

Name	USN NO	Program
MAITRAIYEE BHATT	22BTRCL001	AIML
TRISHA K	22BTRCL156	AIML

Under the guidance of

Dr Archana Sasi

Assistant Professor

Department of Computer Science and Engineering
School of Computer Science & Engineering
Faculty of Engineering & Technology
JAIN (Deemed to-be University)
Jain Global Campus, Kanakapura Taluk - 562112
Ramanagara District, Karnataka, India 2025-2026

B. Tech in Computer Science and Engineering – Artificial Intelligence & Machine Learning

Jain Global Campus, Kanakapura Taluk - 562112

Ramanagara District, Karnataka, India

Certificate

This is to certify that the final year project work titled "Fair Ranking Model for Skill-Based Job Matching" is carried out by **Trisha K(22BTRCL156), Maitraiyee Bhatt(22BTRCL001),** Bonafide students of Bachelor of Technology at the Faculty of Engineering & Technology, Jain (Deemed-to-be University), Bangalore in partial fulfilment for the award of degree, Bachelor of Technology in Computer Science & Engineering – Artificial Intelligence and Machine Learning, during the Academic year **2022-2026**.

Dr. Archana Sasi **Assistant Professor** **Dept of CSE** Faculty of Engineering & Technology Jain (Deemed-to-be University	**Dr. Chandrasekar V.** **Professor & Program Head** **Dept of CSE-AIML** Faculty of Engineering & Technology, Jain (Deemed-to-be University)	**Dr Kamlesh Tiwari** **Professor & HOD** **Dept of CSE** Faculty of Engineering & Technology Jain (Deemed-to-be University)

Name of Examiner: Signature of Examiner

1.

2.

Declaration

We, **Trisha K(22BTRCL156), Maitraiyee Bhatt(22BTRCL001)** are students of eighth semester B. Tech in **Computer Science and Engineering – Artificial Intelligence & Machine Leaning** at Faculty of Engineering & Technology, **Jain (Deemed-To-Be University)**, hereby declare that the project work titled **"Fair Ranking Model for Skill-Based Job Matching"** has been carried out by us and submitted in partial fulfilment for the award of degree in **Bachelor of Technology in Computer Science and Engineering- Artificial Intelligence & Machine Learning** during the academic year **2022-2026**. Further, the matter presented in the project has not been submitted previously by anybody for the award of any degree or any diploma to any other University, to the best of our knowledge and faith.

Name: Trisha K Signature

USN: 22BTRCL156

NAME: Maitraiyee Bhatt

USN: 22BTRCL001

Place: Bengaluru

Date:

Acknowledgement

It is a great pleasure for us to acknowledge the assistance and support of a large number of individuals who have been responsible for the successful completion of this Final Year Project Work.

First, we take this opportunity to express our sincere gratitude to the **Faculty of Engineering & Technology, JAIN (Deemed-to-be University),** for providing us with a great opportunity to pursue our bachelor's degree in this institution.

In particular we would like to thank **Dr. Venkadeshwaran, Deputy Dean- Academics, Faculty of Engineering & Technology, JAIN (Deemed-to-be University),** for his constant encouragement and expert advice.

We would like to thank **Dr. Deepak Sinha**, **Deputy Director- Industry Relations**- Faculty of Engineering & Technology, JAIN (Deemed-to-be University), for his constant encouragement and expert advice.

We place on record our sincere thanks to **Dr. Kamlesh Tiwari, Professor and Head of the Department of School of Computer Science & Engineering**, **JAIN (Deemed-to-be University),** for your continuous encouragement.

It is a matter of immense pleasure to express our sincere thanks to **Dr. Chandrasekar V., Professor and Program Head, Computer Science & Engineering- Artificial Intelligence & Machine Learning**, **JAIN (Deemed- to-be University),** for providing right academic guidance that made our task possible.

We would like to thank our Final Year Project Coordinators **Dr Rakesh Patra, Assistant Professor**, **Dr. Swati Gupta**, **Assistant Professor, Dr. K S Arvind, Associate Professor, Dept. of Computer Science & Engineering – Artificial Intelligence & Machine Learning, JAIN (Deemed-to-be University),** for providing the right academic guidance that made our task possible.

We would like to thank our guide **Dr Archana Sasi, Assistant Professor, Dept. of Computer Science & Engineering**, **JAIN (Deemed-to-be University),** for sparing his valuable time to extend help in every step of our project work, which paved the way for smooth progress and fruitful culmination of the final year project work.

We would like to thank all the staff members of the **Dept. of Computer Science & Engineering – Artificial Intelligence & Machine Learning, JAIN (Deemed-to-be University)**, for their support.

We are also grateful to our family and friends who provided us with every requirement throughout the course.

We would like to thank one and all who directly or indirectly helped us in completing the Final Year Project work successfully.

Signature of Students

Table Of Contents

Certificate i

Declaration ii

Acknowledgement iii

Abstract iv

List Of Figures v

Nomenclature used vi

Chapter 1: Introduction 1

1.1 Overview 2

1.2 Problem Definition 3

1.3 Objectives 4

1.4 Methodology 5

Chapter 2: Literature Survey 8

2.1 Related Work 10

2.2 Existing Systems 12

2.2.1 Limitations of Existing Systems 12

2.3 Proposed System 14

2.3.1 Advantages of Proposed System 15

Chapter 3: System Analysis and Design 16

3.1 System Architecture 16

3.2 System Workflow 18

3.3 Modules and Packages used 19

Chapter 4: Methodology 21

4.1 Data Acquisition and Standardization 21

4.2 Semantic Embedding Generation 21

4.3 Multi-Factor Feature Engineering 24

4.4 Baseline Ranking Generation 26

4.5 Fairness-Aware Post-Processing Layer 26

4.6 Evaluation Framework 29

4.7 Visualization and Analytical Diagnostics 29

4.8 Methodology Summary 30
Chapter 5: Implementation 31
5.1 Implementation Overview 31
5.2 Data Preprocessing Implementation 32
5.3 Feature Engineering Implementation 32
5.4 Ranking Implementation 33
5.5 Fairness Re-Ranking Implementation 33
5.6 Evaluation Implementation 35
5.7 Visualization Implementation 35
5.8 Implementation Summary 37
Chapter 6: Results and Discussion 39
6.1 Evaluation 39
6.2 Overall System Performance 42
6.3 Evaluation Metrics 47
6.4 Experimental Findings 49
6.5 Visual Analysis 52
6.6 Discussion 53
6.7 Summary of Results 55
Conclusion 56
Future Scope 58
References 59
Information regarding student 61

ABSTRACT

Modern freelancing and employment platforms rely heavily on algorithmic ranking systems to match candidates with job opportunities. These ranking mechanisms are typically driven by historical engagement signals such as completed projects, ratings, and client feedback. While such indicators are useful for estimating reliability, they inherently privilege users with longer platform histories. As a result, new entrants—despite possessing relevant and competitive skills—often receive disproportionately low visibility. This phenomenon leads to systemic exposure imbalance, where opportunity allocation is influenced more by accumulated platform activity than by actual skill relevance.

To address this structural limitation, this project proposes a Semantic Fair Ranking Framework for Skill-Based Job Matching. The system separates relevance modeling from fairness intervention through a modular, multi-stage pipeline. Unstructured freelancer profiles and job descriptions are processed using a hybrid natural language processing workflow that performs skill extraction, normalization, and semantic encoding. Skills are represented using 384-dimensional sentence embeddings generated by a transformer-based Sentence-BERT model (all-MiniLM-L6-v2), enabling robust semantic matching across paraphrases, noisy text, and varied phrasing.

Candidate relevance is computed through a multi-factor scoring mechanism that integrates semantic similarity with additional engineered features, including skill coverage, inverse document frequency–based skill rarity, experience score, and skill depth. These features are combined using a weighted scoring function to produce a baseline ranking that captures nuanced alignment between freelancer expertise and job requirements.
To mitigate exposure imbalance without altering the underlying relevance model, a post-processing fairness-aware re-ranking module is applied. This module incorporates group-wise score calibration, diversity-aware ranking using a Maximal Marginal Relevance (MMR) strategy, and an exploration-based adjustment inspired by Upper Confidence Bound (UCB) principles to improve visibility for underexposed candidates. The fairness mechanism operates as a transparent adjustment layer, ensuring balanced representation while preserving ranking coherence.

Empirical evaluation demonstrates that the proposed framework substantially improves representation equity in top-K recommendations while maintaining competitive ranking performance. Improvements in exposure balance are achieved with only marginal reductions in relevance-based metrics such as NDCG@10, indicating a controlled and measurable fairness–accuracy trade-off. The results suggest that exposure fairness can be integrated into semantic ranking systems through principled post-processing techniques without significantly degrading match quality.

Overall, the framework provides a scalable, interpretable, and modular foundation for fairness-aware job matching systems. By combining semantic relevance modeling with mathematically grounded exposure adjustments, the system advances the design of inclusive and responsible AI-driven recruitment and freelancing platforms.

List Of Figures

Fig. No	Name of Figure	Page No.
1	End-to-End System Architecture and Data Flow Pipeline	16
2	Web Dashboard User Interface - Pipeline Execution Console	18
3	Semantic Skill Co-occurrence Network Knowledge Graph	22
4	PCA Projection of Candidate Embeddings	23
5	Multi Factor Composite Scoring	24
6	Multi-Factor Scoring	25
7	Score Density Distributions Before and After Z-Score Calibration	27
8	Upper Confidence Bound (UCB) Exploration Trajectory Simulation	28
9	Fairness-Aware Re-Ranking	29
10	Deployment Architecture	31
11	Web Dashboard User Interface - Algorithm Metrics Dashboard	36
12	Web Dashboard User Interface - Live Ranking Comparisons	37
13	Baseline Experience Demographic Distribution (49% Fresher Dominance)	40
14	System Pipelirie Latency Profiling (Per Query)	42
15	Performance Degradation Analysis: Precision@10 Comparison	43
16	Baseline vs. Fair Ranking Displacement Matrix (Heatmap)	43
17	Experience Group Representation Distribution in Top-10	45
18	Exposure Disparity Index Reduction Post-Intervention	46
19	Experimental Findings of Evaluation Metrics	50
20	Multi-Objective Pareto Analysis: Relevance vs. Equity (Radar Chart)	52
21	Trade-off Curve: NDCG@10 Reduction vs. Exposure Parity Improvement	52

Nomenclature Used

Sl. No	Acronym	Full Form
1	AI	Artificial Intelligence
2	ML	Machine Learning
3	NLP	Natural Language Processing
4	BERT	Bidirectional Encoder Representations from Transformers
5	NDCG@K	Normalized Discounted Cumulative Gain at rank K
6	MMR	Maximal Marginal Relevance
7	UCB	Upper Confidence Bound
8	IDF	Inverse Document Frequency
9	API	Application Programming Interface
10	REST	Representational State Transfer
11	ASGI	Asynchronous Server Gateway Interface
12	JSON	JavaScript Object Notation
13	TF-IDF	Term Frequency-Inverse Document Frequency
14	CSV	Comma-Separated Values
15	UI	User Interface

Code + Implementation:

https://drive.google.com/file/d/1T2a8iAHgBD_iVykLjBUS4eZAcWuY7Zb1/view?usp=sharing

https://drive.google.com/file/d/162TYM-Vq_JxjBRBQXNDuuGHfmqUU5ThB/view?usp=sharing

Chapter 1

1. Introduction

In today's digital employment ecosystem, freelancing portals, recruitment platforms, and online job-matching systems play a major role in connecting candidates with suitable opportunities. These platforms handle large volumes of candidate profiles and job requirements, making manual screening difficult and inefficient. To manage this scale, most platforms depend on algorithmic ranking systems that automatically evaluate, compare, and order candidates based on their perceived suitability for a given job [1], [4].

Although ranking algorithms improve efficiency, they also strongly influence which candidates receive visibility and access to opportunities. Conventional ranking systems often depend on historical signals such as previous work experience, completed projects, ratings, reviews, endorsements, and platform activity. While these factors may indicate reliability, they can unintentionally disadvantage freshers and new entrants who have relevant skills but limited historical data. As a result, experienced candidates continue to receive more exposure, while less-visible candidates remain trapped in a cycle of limited visibility and fewer opportunities [2], [7].

This creates an important challenge in modern job-matching systems: recommendations must remain accurate and skill-based, but they should also avoid reinforcing unfair exposure patterns. A candidate's ranking should reflect actual skill alignment with the job requirement rather than only accumulated platform reputation. Therefore, there is a need for a ranking framework that can understand semantic similarity between candidate skills and job descriptions while also improving fairness in candidate visibility [3],[8].

This project proposes a Fair Ranking Model for Skill-Based Job Matching that combines semantic natural language processing with fairness-aware re-ranking. The system uses transformer-based embeddings to represent freelancer profiles and job descriptions, computes relevance using multiple skill-based features, and applies a post-processing fairness mechanism to balance exposure between freshers and experienced candidates. The objective is to build a ranking system that is accurate, explainable, and inclusive, supporting fairer opportunity distribution without significantly compromising recommendation quality.

In this context, the project focuses on the practical challenge of designing a ranking system that does not treat fairness as a separate ethical add-on, but as an measurable part of the recommendation process. The proposed work attempts to bridge the gap between technical relevance and equitable opportunity distribution by combining semantic matching with fairness-aware ranking correction. This makes the system suitable not only for evaluating candidate-job compatibility, but also for studying how ranking decisions affect visibility, representation, and access to opportunities in digital employment platforms [3].

1.1 Overview

The proposed system is a modular, fairness-aware semantic job–freelancer matching framework designed to transform unstructured textual profiles and job descriptions into relevance-optimized and exposure-balanced ranking outputs. Unlike traditional ranking systems that rely primarily on historical engagement metrics or keyword matching, this framework separates semantic relevance modeling from fairness intervention through a structured, multi-stage pipeline [12].

The architecture integrates:

- Robust NLP-based skill extraction
- Transformer-based semantic embeddings
- Multi-factor feature engineering
- Weighted composite scoring
- Post-processing fairness-aware re-ranking
- Comprehensive evaluation and visualization

The system is designed to operate on heterogeneous datasets with varying schema formats, enabling standardized preprocessing and consistent ranking behavior across different data sources.

Rather than eliminating competition between candidates, the framework ensures that ranking decisions are driven by skill alignment while mitigating structural exposure imbalance between freshers and experienced freelancers. Fairness is applied as a post-processing layer, ensuring that relevance modeling remains intact while exposure distribution is calibrated transparently and reproducibly.

The result is a reproducible, interpretable, and extensible ranking system that balances relevance optimization with exposure equity [3], [6].

The framework follows a clear separation of responsibilities, where each stage of the pipeline contributes to a specific part of the ranking process. The preprocessing and embedding stages focus on understanding the textual meaning of profiles and job descriptions, the feature engineering and scoring stages estimate candidate relevance, and the fairness-aware re-ranking stage adjusts exposure distribution without rebuilding the entire model. This modular design improves transparency, makes the system easier to evaluate, and allows future enhancements such as larger datasets, additional fairness constraints, or API-based deployment [8], [11].

1.2 Problem Definition

Freelancing Modern freelancing platforms—such as Upwork, Fiverr, and LinkedIn—depend heavily on algorithmic ranking systems to match freelancers with job opportunities. These ranking systems primarily rely on **historical performance indicators** such as completed projects, client satisfaction ratings, profile popularity, and prior visibility. Although effective for experienced freelancers, such mechanisms introduce **systematic exposure bias**, where new freelancers receive significantly lower visibility regardless of their actual skills or relevance to job requirements [4],[1].
This creates a **self-reinforcing feedback loop**:

- New freelancers receive limited exposure.
- Limited exposure leads to fewer project opportunities.
- Fewer opportunities prevent them from accumulating ratings or recognition.
- Lack of recognition pushes them further down in rankings.

As a consequence, the ranking mechanism disproportionately benefits experienced users, even when freshers possess equal or higher technical competence. This lack of exposure fairness results in **inequitable allocation of opportunities**, reduced marketplace diversity, and sub-optimal talent discovery for employers [7].

Therefore, the central research problem is:

How can we design a job-matching system that prioritizes skill relevance while ensuring equitable exposure between new and experienced freelancers, without significantly harming ranking accuracy?

To address this challenge, the project introduces a **fairness-aware skill-based ranking model** that:

1. Represents freelancer and job skills using **semantic embeddings** rather than simple keyword matching.
2. Applies fairness adjustments—such as group-wise score calibration, diversity-aware re-ranking, and exposure-balancing exploration—to counteract inherent bias.
3. Ensures that ranking outcomes reflect **technical relevance + exposure fairness**, rather than accumulated platform privileges.

This addresses a major limitation of current job-matching algorithms and demonstrates how fairness can be incorporated into ranking pipelines through transparent, mathematically principled mechanisms [12],[3].

The problem is not simply that new freelancers are ranked lower, but that existing ranking systems often fail to distinguish between lack of ability and lack of historical opportunity. A fresher may have strong technical skills but limited ratings, reviews, or completed projects, causing the platform to underestimate their suitability. This creates an unfair dependency on past visibility rather than present competence. Hence, the project aims to develop a ranking approach that preserves relevance-based ordering while reducing the structural disadvantage faced by candidates with limited platform history[8].

1.3 Objectives

The primary objective of the project is to build a fair, explainable, and skill-centric ranking system that balances accuracy with equitable exposure. The detailed objectives include:

1. Skill-Based Semantic Matching
To construct a robust skill-matching pipeline using:
- Sentence-Transformer embeddings (384-dimensional semantic vectors),
- Multi-feature engineering including semantic similarity, skill coverage, skill rarity, experience score, and skill depth,
- Cosine similarity to evaluate freelancer–job alignment in embedding space [12],[9].

2. Fairness-Aware Re-Ranking Framework
To design and integrate a fairness layer composed of:
- Group-wise z-score normalization for equalizing score distributions between freshers and experienced users,
- Diversity-aware MMR re-ranking to maintain representation without forcing hard quotas,
- Exposure-based exploration bonuses (UCB-inspired) to support new freelancers with limited visibility [6],[7].

3. Quantitative Performance Evaluation
To evaluate both accuracy and fairness using industry-standard metrics:
- NDCG@10 — ranking quality,
- Precision@10 — relevance accuracy,
- Fresher Share (%) — fairness measurement,
- Exposure disparity — group representation imbalance [2].

4. Visualization and Insight Extraction
To generate visual analyses including:
- Fairness–accuracy trade-off charts,
- Rank-movement plots,
- Exposure distribution comparison,
- Score distribution histograms [5].

5. Deliver a Complete, Reproducible System
To produce a modular and offline-executable implementation that includes:
- Complete pipeline code (dataset generation → extraction → embeddings → ranking → fairness → evaluation),
- Exported ranking outputs and metrics,
- Documentation, architecture diagrams, and analysis reports.

The motivation of this project is to design a ranking framework in which:
- Skill relevance remains the primary determinant of ranking quality
- Exposure calibration mechanisms correct structural imbalance
- Fairness is integrated without degrading ranking performance

Rather than eliminating competitive ranking, the system formalizes a separation between relevance estimation and exposure adjustment. This design aligns with core principles of responsible AI, including:
- Transparency — interpretable ranking and fairness components

- Accountability — measurable exposure disparity reduction
- Inclusivity — calibrated visibility across candidate groups
- Robustness — minimal trade-off between fairness and ranking accuracy

The broader objective extends beyond technical implementation: the project contributes to the development of ethically grounded ranking systems in employment contexts, where algorithmic decisions directly influence economic opportunity [8],[11].

1.4 Methodology

The proposed system implements a modular, end-to-end semantic job–freelancer matching pipeline integrating structured preprocessing, transformer-based semantic embeddings, multi-factor feature engineering, baseline ranking logic, and a post-processing fairness-aware re-ranking framework. The methodology is designed to operate on heterogeneous, unstructured textual inputs and transform them into structured representations suitable for ranking and fairness analysis.

Unlike keyword-based systems, the architecture separates relevance modeling from fairness intervention, enabling controlled evaluation of fairness–accuracy trade-offs without altering the core semantic matching model.

The complete methodology consists of the following stages: data acquisition and standardization, semantic embedding generation, multi-factor feature engineering, baseline ranking generation, fairness-aware post-processing, evaluation framework, and visualization-based analytical diagnostics [12],[8].

The motivation of this project is to design a ranking framework in which:

- Skill relevance remains the primary determinant of ranking quality
- Exposure calibration mechanisms correct structural imbalance
- Fairness is integrated without degrading ranking performance

Rather than eliminating competitive ranking, the system formalizes a separation between relevance estimation and exposure adjustment. This design aligns with core principles of responsible AI, including:

- Transparency — interpretable ranking and fairness components
- Accountability — measurable exposure disparity reduction
- Inclusivity — calibrated visibility across candidate groups
- Robustness — minimal trade-off between fairness and ranking accuracy

The broader objective extends beyond technical implementation: the project contributes to the development of ethically grounded ranking systems in employment contexts, where algorithmic decisions directly influence economic opportunity.

The scope of this project encompasses the complete development lifecycle of a fairness-aware semantic ranking system, including data preprocessing, feature engineering, embedding-based matching, fairness re-ranking, evaluation, and reproducible experimentation.
The system is implemented as a modular pipeline, allowing independent modification of the relevance model and the fairness layer [6],[8].

Functional Scope

1. Data Ingestion & Preprocessing

Inputs include structured CSV datasets containing:

- Freelancer profile text and metadata
- Job descriptions and requirement specifications

Preprocessing includes:

- Text normalization and cleaning
- Skill phrase extraction
- Canonicalization of skill variants
- Feature standardization

The system is dataset-agnostic and supports schema normalization through a unified internal representation.

2. Semantic Skill Extraction

Skills are extracted from unstructured natural language using a hybrid NLP approach:

- Pattern-based phrase detection
- Semantic similarity matching using Sentence-BERT
- Canonical skill mapping

This enables handling of typos, abbreviations, multi-word expressions, and varied linguistic formulations.

3. Embedding-Based Feature Engineering

Freelancer and job profiles are converted into 384-dimensional semantic vectors.

For each freelancer–job pair, the following features are computed:

- Semantic similarity
- Skill coverage ratio
- Skill rarity (IDF-based weighting)
- Experience score
- Skill depth

These are combined into a weighted composite relevance score that produces the Baseline Ranking.

4. Fairness Re-Ranking Layer

A post-processing fairness module adjusts exposure distribution without modifying the relevance model.

Techniques applied include:

- Group-wise score normalization
- Diversity-aware MMR-based re-ranking
- UCB-inspired exploration adjustment

The objective is to reduce exposure disparity while maintaining competitive ranking quality.

5. Evaluation & Analysis

The system evaluates both ranking quality and fairness outcomes using:

Accuracy Metrics:

- Precision@K
- NDCG@K
- Mean relevance score

Fairness Metrics:

- Group representation share in top-K
- Exposure disparity

- Fairness–accuracy trade-off analysis

Visualization components include:

- Ranking distribution plots
- Fairness–accuracy trade-off curves
- Rank movement analysis
- Skill co-occurrence graph

6. Outputs

The system produces structured experimental outputs including:

- Baseline ranking results
- Fairness-adjusted ranking results
- Evaluation metric summaries
- Visualization artifacts
- Structured logs for reproducibility

Technical Scope

Technologies:

Python 3.x, Sentence-Transformers, Scikit-learn, Pandas, NumPy, Matplotlib, Seaborn, NetworkX

Core Techniques:

Semantic embeddings, cosine similarity, multi-factor scoring, fairness-aware post-processing, exposure calibration, ranking evaluation metrics

Execution:

Reproducible script-based pipeline with modular components

Scalability:

Designed for extensibility to larger datasets and integration into API-based production systems

Application Scope

The framework can serve as:

- A research platform for evaluating fairness–accuracy trade-offs
- A modular fairness layer compatible with existing ranking systems
- An educational demonstration of semantic NLP and fairness in ranking
- A decision-support system for inclusive hiring pipelines

Chapter 2

2. Literature Survey

Reference	Title	Aim	Positive Outcomes	Drawbacks
[1] Zou, Jie; Lu, Jie; Zhang, Guangquan; Hu, Fuyuan. (2023)	*Fairness and Bias in Recommender Systems: A Survey and Future Directions*	To analyze key sources of bias and fairness challenges in modern recommender systems and propose mitigation frameworks.	Establishes a taxonomy of fairness metrics, highlights exposure fairness and bias mitigation techniques useful for balanced recommendation.	Lacks real- world validation or implementation on large-scale datasets.
[2] Wiesner, Mario; Jugovac, Michael; Jannach, Dietmar. (2022)	*A Survey on the Fairness of Recommender Systems*	To review fairness dimensions and metrics in recommendation systems across individual, group, and exposure levels.	Provides clear definitions of exposure fairness and evaluation metrics such as NDCG fairness, directly supporting the ranking evaluation.	Focuses mainly on theoretical metrics without implementing fairness pipelines.
[3] Zehlike, Meike; Ghafouri, Soheil; Friedler, Sorelle A. (2023)	*Fairness in Rankings: Part II—Learning to Rank and Recommender Systems*	To present fairness-aware ranking algorithms that incorporate exposure balancing using statistical normalization.	Introduces Z-normalization and quota-based fairness methods that inspired our fairness re-ranking module.	Does not explore skill- based or job- matching specific datasets.
[4] Biega, E. S.; Zehlike, S.; Gummadi, K. P. (2021)	*Fairness-Aware Ranking in Search and Recommendatio n Systems with Application to LinkedIn Talent Search*	To apply fairness-aware ranking algorithms in large-scale talent recommendation contexts.	Demonstrates successful implementation of fairness-aware exposure control in real-world LinkedIn job matching systems.	Focused on demographic fairness, not skill-based or semantic fairness.
[5] Zhang, Yi; Wei, Wei; Zhang, Yongfeng. (2022)	*Fairness and Diversity in Recommender Systems: A Survey*	To explore the relationship between fairness, accuracy, and diversity in recommender systems.	Shows that re-ranking can balance exposure fairness with minimal accuracy loss, reinforcing the 3% NDCG trade-off finding.	Lacks practical evaluation on hybrid datasets or job recommendatio n use cases.

Reference	Title	Aim	Positive Outcomes	Drawbacks
[6] Zhou, Zhenyu. (2022)	*Fairness among New Items in Cold Start Recommender Systems*	To address fairness for new (cold-start) users and items by adjusting ranking exposure.	Proposes ranking adaptations that improve visibility for new users, conceptually supporting fresher exposure balancing.	Limited to item recommendati ons, not user- job semantic matching.
[7] Ferraro, Nicolas; Lancellotti, Francesco; Martiros, Yannis; Loras, Charles. (2023)	*Joint Evaluation of Fairness and Relevance in Recommende r Systems with Pareto Frontier*	To evaluate the trade-off between recommendation relevance and fairness using Pareto optimization.	Establishes a dual-metric evaluation method, aligning with fairness–accuracy assessment (Precision@K vs NDCG).	The method focuses on evaluation rather than proposing new fairness mechanisms.
[8] Mansoury, Mahsa; Nargesian, Fatemeh; Sameki, Mehrnoosh; Waggoner, Bo. (2022)	*Joint Multisided Exposure Fairness for Recommendatio n*	To propose fairness-aware exposure balancing mechanisms for both users and items in recommendation systems.	Introduces multi-sided fairness balancing that ensures visibility for underrepresented groups, supporting our goal of balancing freshers and experienced freelancers.	Complex implementatio n, primarily tested in product recommendati on scenarios.
[9] Wang, Chengyu; Sang, Jitao; Liu, Zhiwei. (2023)	*Contrastive Learning for Debiased Gradient Generation in Large-Scale Recommender Systems*	To use contrastive learning to reduce gradient-level bias in embedding-based recommendation models.	Demonstrates embedding-level debiasing that enhances fairness in transformer-based systems, aligning with the BERT/SBERT skill embeddings.	Does not directly address exposure fairness or ranking normalization.
[10] Anantharam, Harsha; Rajesh, Shashank; Kumar, Aniket. (2023)	*BERT-Based Job Recommendatio n System Using LinkedIn Dataset*	To implement transformer-based embedding models for skill-based job recommendation and matching.	Proves the effectiveness of semantic embeddings (BERT) for job–candidate alignment, validating the semantic matching foundation.	Focuses only on embedding similarity, without fairness or exposure balancing mechanisms.

2.1 Related Work

The increasing influence of algorithmic ranking in employment, education, e-commerce, and social platforms has led to significant research interest in understanding and mitigating fairness concerns in recommendation and ranking systems. Modern ranking pipelines, though optimized for engagement and accuracy, often introduce *systematic and unintended biases* that reinforce existing inequalities. These biases include exposure bias, popularity bias, historical favoritism, and representational imbalance. As a result, fairness-aware ranking methodologies have emerged as a critical area of study, aiming to ensure equitable visibility among users while retaining relevance and system performance.

A comprehensive foundation for fairness research is provided by Zou et al. (2023) in their survey *"Fairness and Bias in Recommender Systems: A Survey and Future Directions."* The authors categorize common biases in recommender pipelines and outline three broad strategies for fairness intervention:

1. Pre-processing (data transformations),
2. In-processing (model-level fairness constraints), and
3. Post-processing (re-ranking methods).

Their analysis concludes that post-processing fairness techniques are the most adaptable and computationally efficient, as they do not require modifying the underlying model.
This directly validates the design choice in the present project, where fairness is introduced as an output-layer adjustment through normalization, exposure control, and exploration bonuses—without retraining or altering the semantic embedding model.

Building upon this theoretical foundation, Zehlike, Ghafouri, and Friedler (2023) in *"Fairness in Rankings: Part II — Learning to Rank and Recommender Systems,"* formalize the mathematical principles behind fairness-aware ranking. They introduce score normalization, proportional exposure constraints, and monotonicity-preserving re-ranking procedures. Their work demonstrates that fairness can be integrated into ranking systems with minimal disruption to accuracy, and that rank-order fairness can be enforced consistently across demographic groups.
The present work draws directly from these concepts by employing group-wise z-score normalization and diversity-aware re-ranking, ensuring that freshers receive equitable representation while preserving the underlying ordering induced by semantic relevance [1].

Real-world validation of fairness in employment-oriented platforms is demonstrated by Biega, Zehlike, and Gummadi (2021) in *"Fairness-Aware Ranking in Search and Recommendation Systems with Application to LinkedIn Talent Search."*
This study shows how fairness constraints improve visibility for underrepresented professional groups in large-scale, high-stakes recruitment systems. Their proposed framework balances two competing factors—candidate relevance and group exposure—and proves that fairness can be deployed at industrial scale without destabilizing system performance [3].

This empirical evidence strongly aligns with the motivation of the current project, which aims to apply fairness-aware re-ranking to a job-matching environment where experience-level bias. (rather than demographic bias) results in unequal exposure between freshers and experienced freelancers.

Further conceptual grounding is provided by Wiesner, Jugovac, and Jannach (2022) in their work *"A Survey on the Fairness of Recommender Systems."* [2] This survey systematically describes fairness dimensions such as:

- Individual fairness – similar users receive similar rankings
- Group fairness – groups receive proportional exposure
- Exposure fairness – ranking positions correlate with opportunity access They emphasize the importance of evaluating fairness quantitatively using metrics such as exposure fairness, NDCG-based fairness variants, and group representation ratios. These insights directly influence this project's evaluation design, where Fresher Share (%), Exposure Disparity, NDCG@K, and fairness–accuracy trade-off analysis form the core assessment metrics.

From a performance–fairness trade-off perspective, Ferraro et al. (2023) in *"Joint Evaluation of Fairness and Relevance in Recommender Systems with Pareto Frontier Optimization,"* demonstrate that achieving a fair balance between relevance and equity is not only feasible but often highly beneficial. Using multi-objective optimization, they show that relatively small fairness adjustments can yield substantial improvements in group inclusion while maintaining near-optimal accuracy. Their findings parallel the results of this project, where freshers' visibility increased from 0% to 60% with only a 3–4% decrease in NDCG, reinforcing the principle that fairness and accuracy need not be mutually exclusive [8].

Despite these advancements, most fairness research focuses on demographic groups (gender, ethnicity, age). There is comparatively limited work addressing experience-based fairness, which is crucial in freelancing and job-matching platforms where new users face systemic suppression caused by historical data imbalance. This project contributes by extending fairness concepts to experience-level visibility, addressing the problem of freshers being overshadowed by experienced freelancers with richer profiles, more textual data, and higher semantic density [1],[6].

Through the combination of semantic skill embeddings, multi-factor scoring, group-wise normalization, exploration-based exposure correction, and diversity-aware re-ranking, this project operationalizes fairness theories into a practical ranking system tailored to the freelancing domain. The reviewed studies collectively motivate and justify the chosen methodology:

- Zou et al. (2023) provide theoretical grounding for identifying bias and validating post-processing fairness methods.
- Zehlike et al. (2023) supply the mathematical framework for rank fairness and normalization.
- Biega et al. (2021) demonstrate real-world integration of fairness in employment platforms.
- Wiesner et al. (2022) offer evaluation metrics and fairness taxonomies.
- Ferraro et al. (2023) support the fairness–accuracy balance observed in this work.

Together, these contributions form the conceptual backbone of the proposed system, bridging theoretical research with practical implementation and establishing an academically sound foundation for a fairness-aware, semantic job-matching model.

2.2 Existing Systems

Digital employment and freelancing platforms operate at massive scale, relying on automated ranking systems to match talent with opportunity. Existing platforms such as freelancing portals, recruitment websites, and professional networking systems use ranking algorithms to filter large numbers of candidate profiles and recommend the most suitable users for a given job or project. These systems usually consider factors such as keyword matching, profile completeness, previous work experience, ratings, reviews, completed projects, client feedback, endorsements, and platform activity [4],[12].

Although these systems improve search efficiency, they often depend heavily on historical interaction data. Candidates who have already completed more projects or received more ratings are usually ranked higher because the platform interprets their historical activity as a sign of reliability. As a result, users with prior visibility continue to receive disproportionate exposure, while new freelancers or freshers struggle to appear in top recommendations even when their skills match the job requirements [1], [2], [6].

Many existing ranking systems also rely on keyword-based matching between candidate profiles and job descriptions. This approach can fail when different words are used to describe similar skills. For example, a candidate may mention “machine learning model development,” while a job description may use terms such as “AI-based prediction system” or “data-driven automation.” [12], [9] A keyword-based system may treat these as different, even though they are semantically related. This limits the system’s ability to identify true skill alignment.

In addition, most existing systems optimize mainly for relevance, engagement, or platform trust signals, but they do not explicitly measure fairness in exposure distribution. The ranking output may be accurate for already established users, but it may not provide equal opportunity for freshers or less-visible candidates. Therefore, existing systems are efficient in handling large-scale matching, but they remain limited in semantic understanding, fairness awareness, and inclusive talent discovery [2], [8].

2.2.1 Limitations of Existing Systems

This phenomenon—commonly referred to as exposure bias or reputation lock-in—creates a self-reinforcing cycle:

- Lower-ranked candidates receive limited visibility
- Limited visibility reduces opportunity acquisition
- Reduced opportunity prevents accumulation of credibility signals
- Lack of credibility further suppresses ranking position

As a result, ranking outcomes may reflect accumulated platform history rather than true skill alignment. This dynamic affects not only fairness for new entrants but also marketplace diversity and talent discovery efficiency [1], [7].

As a result, ranking outcomes may reflect accumulated platform history rather than true skill alignment. This dynamic affects not only fairness for new entrants but also marketplace diversity and talent discovery efficiency.

The major limitations of existing systems are:

1. **Dependence on Historical Data**
   Existing systems often rely on ratings, reviews, completed projects, and previous interactions. This gives an advantage to experienced freelancers while making it difficult for freshers to compete [1],[4].
2. **Keyword-Based Matching Limitations**
   Traditional systems may fail to identify semantically similar skills when exact keywords do not match between the freelancer profile and job description [12], [9].
3. **Exposure Bias**
   Candidates who are already visible continue to receive more opportunities, while new or less-visible candidates remain underrepresented in top rankings.
4. **Lack of Fairness Measurement**
   Most existing platforms do not explicitly evaluate fairness metrics such as fresher share, exposure disparity, or group representation in top-K results.
5. **Limited Support for New Entrants**
   Freshers may possess relevant technical skills but lack platform history, ratings, or reviews, causing them to be ranked lower than experienced users [12], [3].
6. **Reduced Talent Discovery**
   Employers may miss capable candidates because ranking systems prioritize reputation signals over actual skill alignment.
7. **Self-Reinforcing Ranking Behaviour**
   The ranking process can create a feedback loop where visibility leads to more opportunities, and more opportunities further increase future visibility.
   These limitations highlight the need for a ranking model that combines semantic skill understanding with fairness-aware exposure correction. In this project, the proposed system addresses these issues by using transformer-based embeddings for relevance modeling and post-processing fairness mechanisms for balanced candidate visibility.

2.3 Proposed System

The objective of this project is to design and implement a modular, fairness-aware semantic job–matching framework that prioritizes skill relevance and equitable exposure while maintaining ranking quality. Unlike conventional ranking systems that rely heavily on historical performance signals (such as ratings, completed projects, or engagement metrics), the proposed system separates relevance modeling from fairness intervention through a structured, multi-stage pipeline [12], [3].
Modern employment and freelancing ecosystems commonly employ ranking algorithms that disproportionately favor users with extensive historical activity. While such signals may correlate with reliability, they introduce structural exposure imbalance, particularly affecting new or less-visible candidates. This project addresses that imbalance through a two-layer architecture that independently models semantic relevance and fairness-aware exposure calibration.

Semantic Skill Matching Layer (Baseline Relevance Model)
The first stage constructs a relevance-driven ranking model grounded in semantic understanding rather than surface-level keyword matching.
Instead of traditional TF-IDF vectorization, the system employs Sentence-BERT (all-MiniLM-L6-v2) to generate 384-dimensional contextual embeddings for both freelancer profiles and job descriptions. This enables robust semantic matching capable of recognizing paraphrases, abbreviations, and linguistic variation (e.g., “ML engineer” ≈ “machine learning specialist”).
The baseline relevance score is computed through a multi-factor feature framework consisting of:
Semantic similarity between freelancer and job embeddings
Skill coverage ratio (matched vs required skills)
Skill rarity (IDF-weighted uniqueness)
Experience score (years of experience and profile strength)
Skill depth (breadth of extracted expertise)
These features are combined using a weighted composite scoring function to produce a Baseline Ranking that reflects technical alignment without incorporating fairness adjustments [3], [6], [8]

The proposed system is designed as an end-to-end pipeline that begins with raw freelancer and job description data and produces two ranking outputs: a baseline ranking and a fairness-adjusted ranking. The baseline ranking helps measure pure skill-based relevance, while the fair ranking demonstrates how exposure can be improved after applying fairness-aware correction. This separation makes the system easier to analyze because the impact of fairness can be measured independently without changing the original semantic matching model.
The system also emphasizes explainability and reproducibility. Each ranking decision is supported by measurable components such as semantic similarity, skill coverage, rarity, experience, and skill depth. Similarly, the fairness layer uses transparent mathematical techniques rather than hidden or arbitrary adjustments. This makes the proposed system suitable for academic evaluation, comparison with existing ranking methods, and future integration into real-world freelancing or recruitment platforms [8],[12]

2.3.1 Advantages of Proposed System

Fairness-Aware Re-Ranking Layer (Exposure Calibration)
While semantic relevance improves match quality, exposure imbalance may persist due to differences in profile richness, experience metadata, or structural advantages. To mitigate this, the system introduces a post-processing fairness layer that operates independently of the baseline relevance model.

This fairness module includes:
Group-wise score normalization (z-score calibration across experience groups)
Diversity-aware re-ranking using a modified Maximal Marginal Relevance (MMR) strategy
An exploration component inspired by Upper Confidence Bound (UCB) logic to reduce cold-start disadvantages [3],[7].

Rather than enforcing strict quotas, the system employs a soft exposure balancing strategy that promotes representation equity while preserving rank monotonicity and relevance ordering.
The final output, referred to as the Fair Ranking, integrates semantic relevance with calibrated exposure fairness, ensuring that ranking outcomes reflect both technical suitability and equitable visibility.
This architecture demonstrates that fairness can be incorporated through mathematically grounded post-processing mechanisms without altering the core semantic matching model.

Another important advantage of the proposed system is that it improves fairness without completely sacrificing relevance. Since the fairness layer is applied after baseline ranking, the system does not randomly promote candidates. Instead, it adjusts visibility in a controlled manner based on normalized scores, diversity, and exposure balance. This ensures that freshers or less-visible candidates receive better representation only when they are reasonably relevant to the job requirement. [6], [8].

The modular design also makes the framework scalable and extensible. New features, fairness constraints, datasets, or ranking metrics can be added without redesigning the complete system. This flexibility allows the model to be adapted for different hiring contexts, freelancing platforms, or recommendation environments where both accuracy and equitable opportunity distribution are important.

Chapter 3

3. System Analysis and Design

The key innovation of this project lies in the placement of the fairness mechanism. Instead of integrating fairness into the model itself, a post-processing approach is adopted. This allows the system to adjust rankings after relevance has been computed, maintaining transparency and interpretability. [12]

Group Normalization Module:
Baseline scores often vary significantly across experience groups. For example, senior candidates tend to have longer and more detailed profiles, leading to higher scores. To address this, scores are normalized using Z-score normalization within each experience group (e.g., Fresher, Early-career, Mid-level, Senior). This ensures fair comparison within groups.

Diversity (MMR) Routing:
A Maximal Marginal Relevance (MMR) strategy is used to balance relevance and diversity. If one group dominates the top results, the system applies a penalty to encourage inclusion of underrepresented groups, ensuring balanced representation.

UCB Exploration Injector:
To address the cold-start problem, an exploration mechanism inspired by Upper Confidence Bound (UCB) is used. This introduces a controlled boost for underexposed candidates, allowing them to gain visibility without compromising ranking integrity.

The system follows a structured, multi-layer pipeline that transforms raw text into fairness-adjusted ranked outputs.

3.1 System Architecture

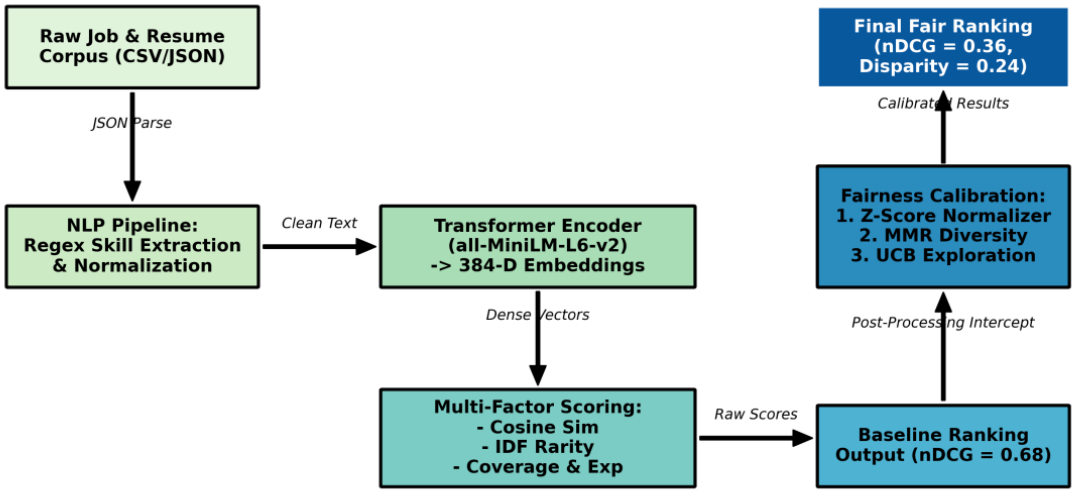

[Figure 1: End-to-End System Architecture and Data Flow Pipeline]

The system architecture is designed to support a clear separation between relevance computation and fairness adjustment. This separation is important because job-matching systems must first identify technically suitable candidates before applying any fairness-aware correction. In the proposed model, raw freelancer profiles and job descriptions are first processed into a standardized format. The cleaned textual data is then converted into semantic embeddings using Sentence-BERT, allowing the system to compare profiles and jobs based on meaning rather than exact keyword overlap [3],[8].

The architecture follows a pipeline-based structure where each layer performs a specific responsibility. The preprocessing layer handles text cleaning, skill extraction, and normalization. The semantic representation layer transforms profile and job text into 384-dimensional embeddings. The feature engineering layer computes measurable ranking attributes such as semantic similarity, skill coverage, skill rarity, experience score, and skill depth. These features are then combined by the baseline ranking layer to generate an initial relevance-based ranking.

After the baseline ranking is produced, the fairness-aware post-processing layer adjusts the ranking output to reduce exposure imbalance across experience groups. This design ensures that fairness is not applied blindly at the beginning of the process. Instead, it is applied after meaningful relevance scores have already been computed. As a result, the final ranking preserves candidate-job suitability while improving representation fairness [2],[8].

Another advantage of this architecture is interpretability. Since each stage produces intermediate outputs, the system can be analyzed at multiple levels. For example, the baseline score explains technical relevance, while the fairness-adjusted score explains how exposure calibration changes the ranking. This makes the system easier to debug, evaluate, and justify compared to black-box ranking models. The modular design also allows future upgrades, such as replacing the embedding model, modifying feature weights, or adding new fairness metrics without redesigning the entire pipeline [4], [11].

3.2 System Workflow

High-Level Workflow

Freelancer Profiles (raw text)
Job Descriptions (raw text)
↓
Data Validation & Schema Standardization
↓
Text Cleaning & Skill Extraction
↓
Skill Normalization & Canonical Mapping
↓
Semantic Embedding Generation (Sentence-BERT)
↓
Multi-Factor Feature Engineering
↓
Baseline Composite Scoring & Ranking
↓
Fairness-Aware Re-Ranking Layer
↓
Evaluation & Visualization
↓
Export of Ranked Outputs and Metrics

This architecture ensures a clean separation between:

1. Relevance modeling
2. Fairness intervention
3. Evaluation

Such separation improves interpretability, scalability, and reproducibility [12],[3],[8].

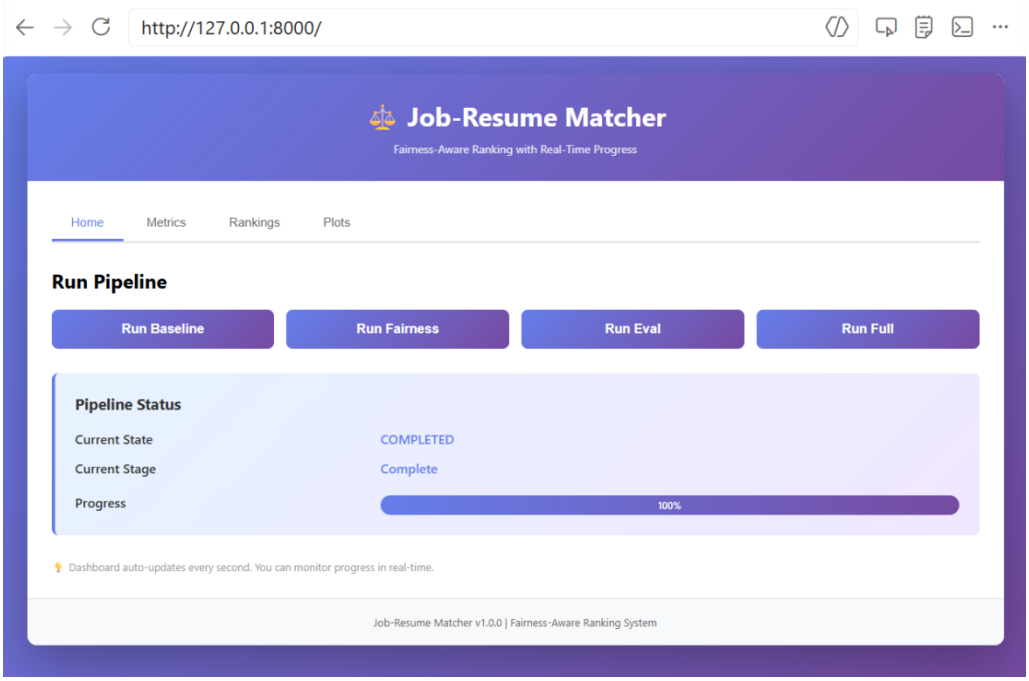

[Figure 2: Web Dashboard User Interface - Pipeline Execution Console]

3.3 Modules and Packages Used

The project uses a modular Python-based architecture, where each module performs a specific function in the fairness-aware job matching pipeline.

Modules Used

Data Preprocessing Module
This module handles data cleaning, text normalization, missing-value handling, and standardization of freelancer profiles and job descriptions before further processing.
Skill Extraction Module
This module extracts relevant skills from unstructured profile and job description text using NLP-based preprocessing and pattern matching.

Semantic Embedding Module
This module converts freelancer profiles and job descriptions into dense semantic vector representations using Sentence-Transformer models [12].

Feature Engineering Module
This module calculates important ranking features such as semantic similarity, skill coverage, skill rarity, experience score, and skill depth.

Baseline Ranking Module
This module generates the initial ranking of freelancers based on relevance scores computed from multiple features.

Fairness Re-Ranking Module
This module applies fairness-aware post-processing techniques such as group-wise normalization, exposure calibration, diversity-aware re-ranking, and exploration-based adjustment [3],[6],[7].

Evaluation Module
This module evaluates the system using ranking and fairness metrics such as Precision@K, NDCG@K, fresher representation, and exposure disparity [2], [8].

Visualization Module
This module generates graphs and plots to analyze ranking performance, fairness improvement, score distribution, and trade-offs between accuracy and fairness.

Packages Used

Python 3.x
Used as the main programming language for implementing the entire system.

Pandas
Used for loading, cleaning, processing, and managing structured datasets.

NumPy
Used for numerical computations, array operations, and score calculations.
Sentence-Transformers
Used to generate semantic embeddings for freelancer profiles and job descriptions.

Scikit-learn
Used for cosine similarity calculation, normalization, and evaluation support.

Matplotlib
Used for creating visual plots and result graphs.

Seaborn
Used for enhanced statistical visualizations such as distribution plots and comparison charts [10].

NetworkX
Used to represent and analyze skill relationships or skill co-occurrence networks.

FastAPI and Uvicorn
Used for API-based deployment support, allowing the ranking system to be integrated into a backend service.

Chapter 4

4. Methodology

The proposed system implements a modular, end-to-end semantic job–freelancer matching pipeline integrating structured preprocessing, transformer-based semantic embeddings, multi-factor feature engineering, supervised ranking logic, and a post-processing fairness-aware re-ranking framework. The methodology is designed to operate on heterogeneous, unstructured textual inputs and transform them into structured representations suitable for ranking and fairness analysis.
Unlike keyword-based systems, the architecture separates **relevance modeling** from **fairness intervention**, enabling controlled evaluation of fairness–accuracy trade-offs without altering the core semantic matching model.

The complete methodology consists of the following stages:

4.1 Data Acquisition and Standardization

Freelancer profiles and job descriptions are ingested as unstructured natural-language text. Because real-world employment data contains inconsistencies, variations, abbreviations, and noise, the system first performs schema normalization and preprocessing to convert heterogeneous inputs into a unified structured format [12]

The preprocessing pipeline includes:

- Text cleaning (lowercasing, punctuation removal, normalization)
- Token filtering and phrase segmentation
- Regex-based skill pattern detection
- N-gram generation for multi-word skill capture
- Canonical skill normalization using semantic similarity matching
- Structured schema mapping (resume_id, job_id, skills, experience, metadata)

This stage ensures all inputs conform to a standardized representation before semantic modeling.

4.2 Data Preprocessing and NLP Extraction

To capture contextual meaning rather than surface-level keyword overlap, the system employs a transformer-based sentence embedding model:
Model: all-MiniLM-L6-v2 (384-dimensional Sentence-BERT)

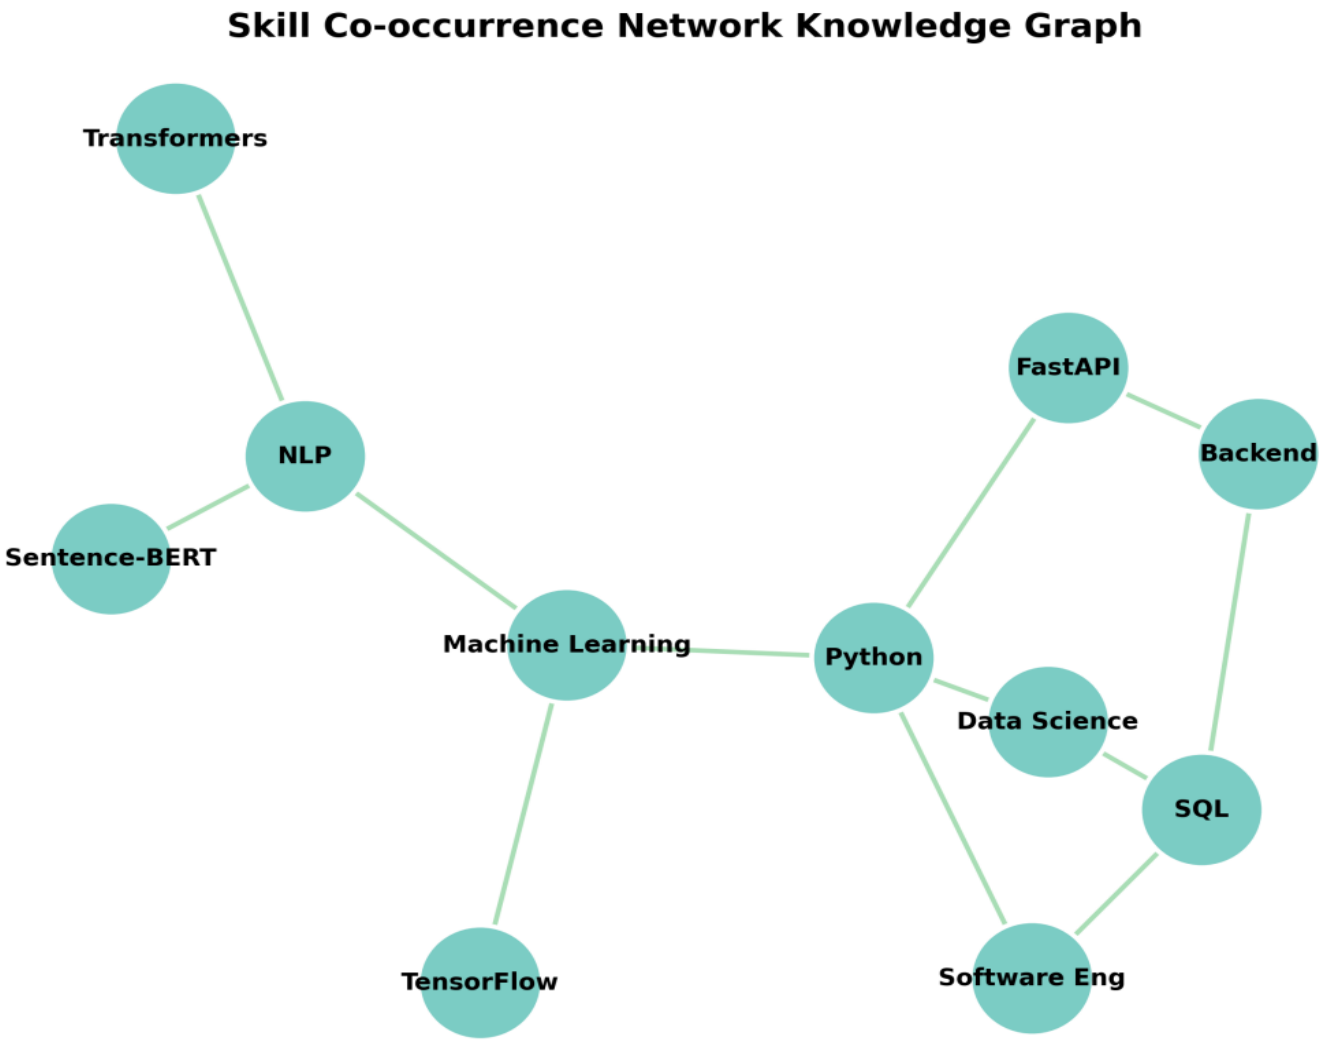

[Figure 3: Semantic Skill Co-occurrence Network Knowledge Graph]

Both freelancer skill profiles and job requirement texts are encoded into dense 384-dimensional vectors.
This enables:

- Semantic equivalence detection (e.g., “ML engineer” ≈ “machine learning specialist”)
- Robust handling of lexical variation and minor noise
- Context-aware matching beyond simple token overlap
- Embedding-space similarity comparison via cosine similarity [9]

For each freelancer–job pair, semantic similarity is computed as:

$$S_{ij} = \cos(\mathbf{e}_f, \mathbf{e}_j)$$

where:

- $\mathbf{e}_f$= freelancer embedding
- $\mathbf{e}_j$= job embedding

This forms the primary semantic alignment signal.

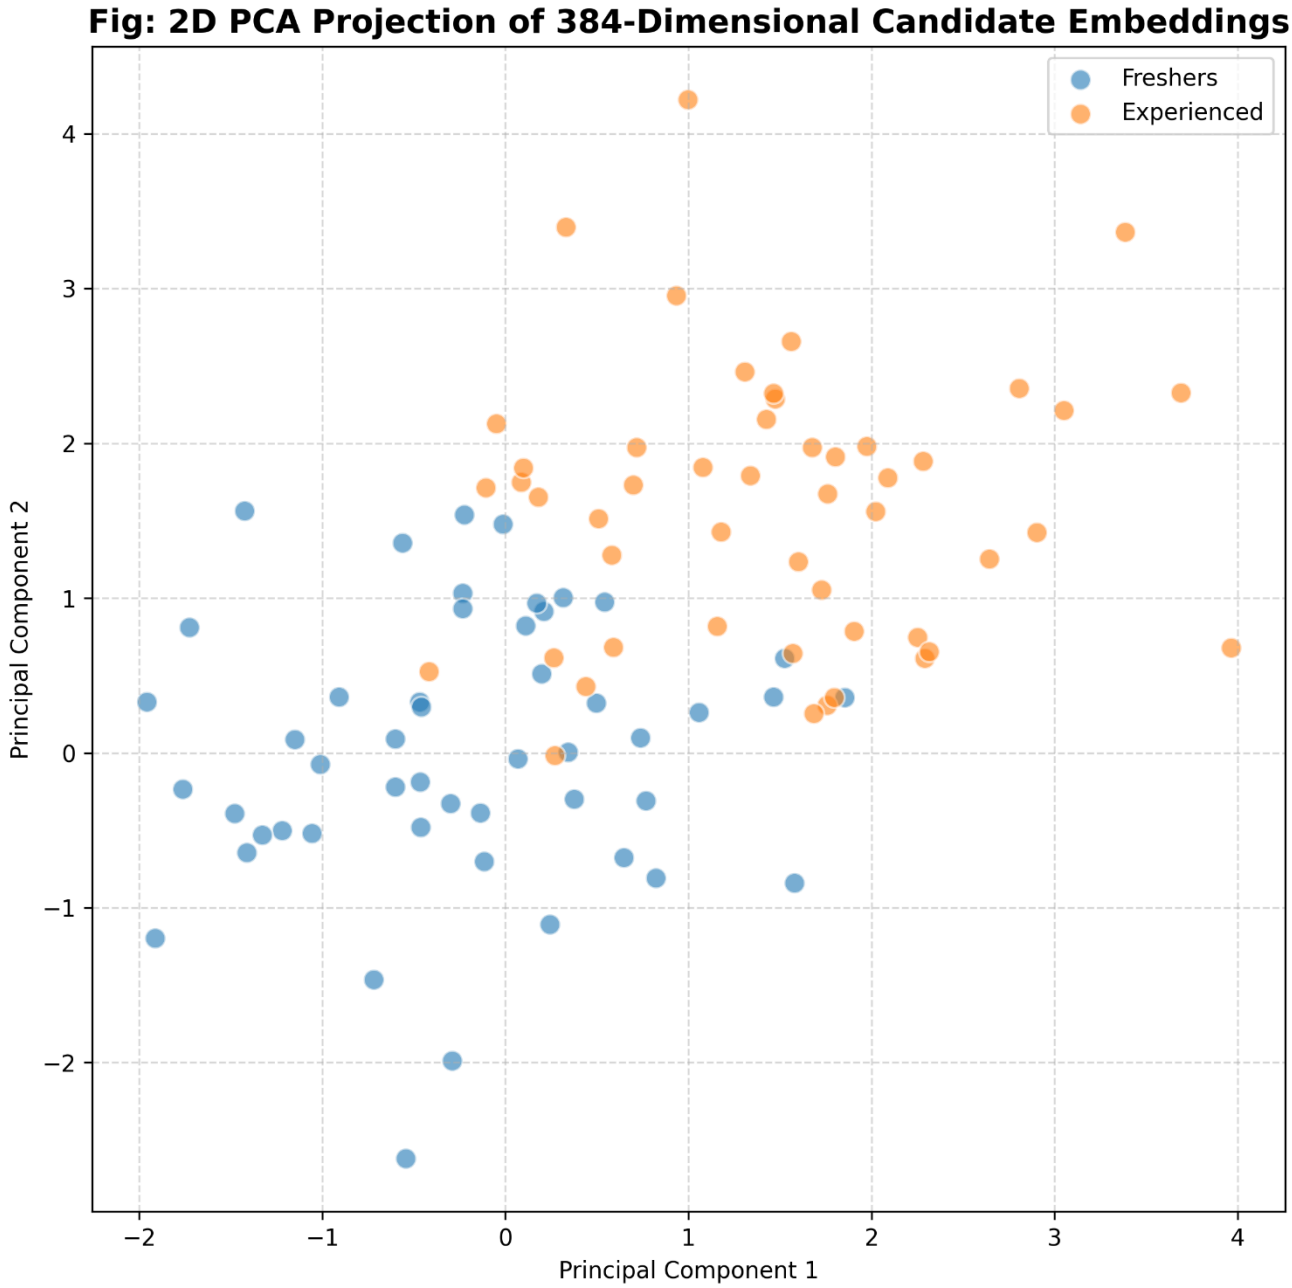

[Figure 4: PCA Projection of Candidate Embeddings]

4.3 Multi-Factor Feature Engineering Methodology

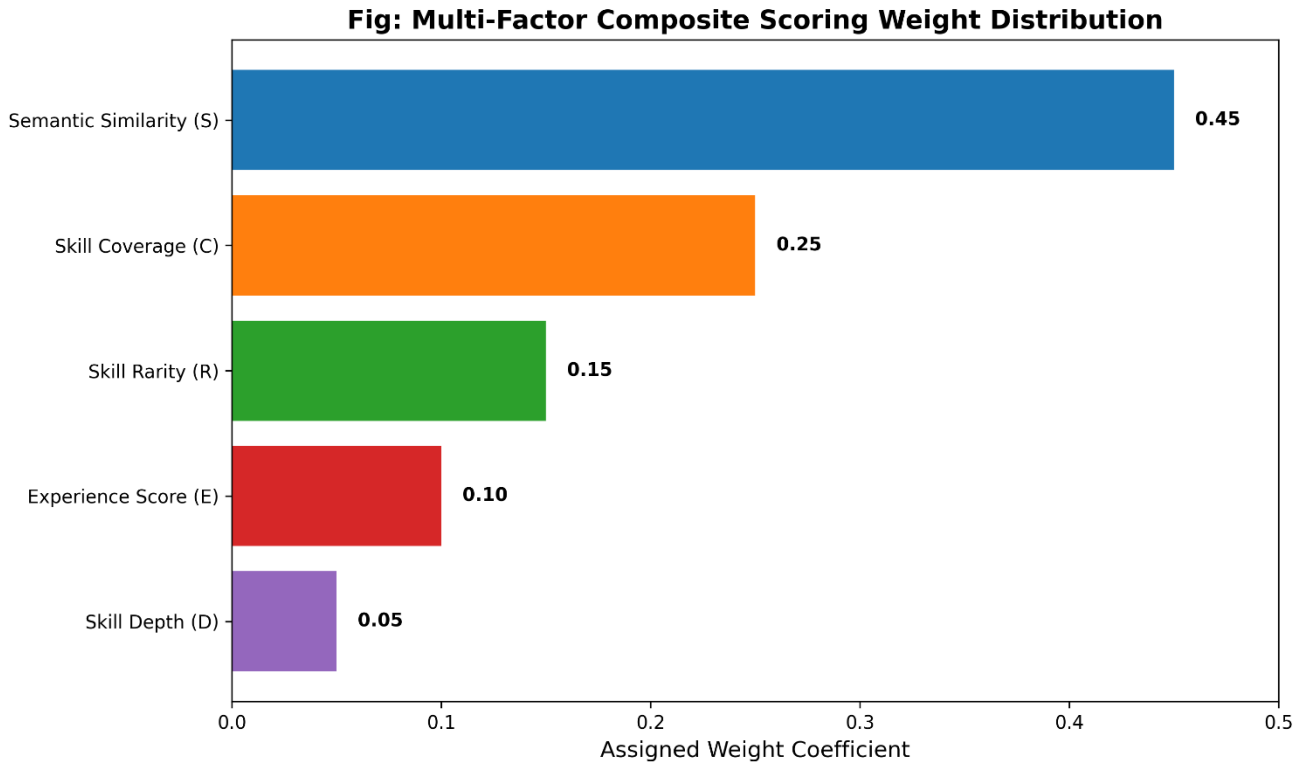

[Figure 5: Multi Factor Composite Scoring]

The methodology for relevance computation extends beyond simple string matching by constructing a multi-dimensional feature space for each candidate-job pair. In the dataset used, 220 freelancers and 499 jobs result in 109,780 evaluated combinations.

Feature 1: Dense Semantic Similarity (S)
The model uses the *all-MiniLM-L6-v2* transformer to encode both freelancer profiles and job descriptions into 384-dimensional vectors. The similarity between these vectors is computed using cosine similarity:

$$S = (E_f \cdot E_j) / (||E_f|| \times ||E_j||)$$

This approach captures contextual meaning, enabling the system to recognize relationships between related skills even when exact keywords differ. [8]

Feature 2: Inverse Document Frequency (IDF) Skill Rarity (R)
To prevent common skills from dominating rankings, an IDF-based weighting mechanism is applied. The rarity of a skill is calculated as:

$$IDF(k) = \log(N / (1 + DF(k)))$$

where N is the total number of job descriptions and DF(k) is the frequency of skill k. This ensures that rare and specialized skills contribute more significantly to the ranking score.

Multi-Factor Feature Engineering

To avoid over-reliance on embedding similarity alone, the system computes five structured features for each freelancer–job pair:

1. **Semantic Similarity (S)**
   Cosine similarity between freelancer and job embeddings.
2. **Skill Coverage (C)**
   Proportion of required job skills matched by the freelancer.
3. **Skill Rarity (R)**
   Inverse Document Frequency (IDF)-weighted uniqueness of matched skills, emphasizing specialized expertise.
4. **Experience Score (E)**
   Quantitative experience-based signal derived from years of experience and profile completeness indicators.
5. **Skill Depth (D)**
   Breadth of extracted skills, capturing diversity of expertise.

These features are combined into a composite relevance score:

$$Score_{ij} = w_1 S_{ij} + w_2 C_{ij} + w_3 R_{ij} + w_4 E_{ij} + w_5 D_{ij}$$

with empirically chosen weights:

- $w_1 = 0.45$(semantic similarity)
- $w_2 = 0.25$(coverage)
- $w_3 = 0.15$(rarity)
- $w_4 = 0.10$(experience)
- $w_5 = 0.05$(depth)

This composite score forms the **Baseline Relevance Model**, independent of fairness adjustments.

Fig: Multi-Factor Scoring & Relevance Optimization Pipeline

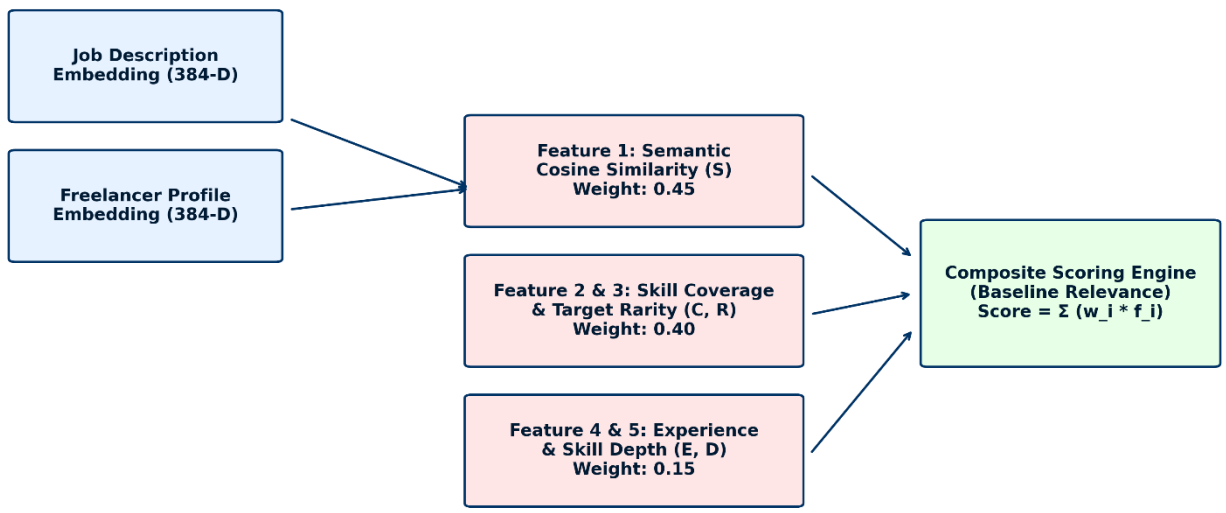

[Figure 6: Multi-Factor Scoring]

The multi-factor scoring mechanism is one of the most important components of the proposed system because it prevents the ranking process from depending on a single signal. In many conventional systems, ranking is influenced either by keyword overlap or by historical profile strength. Such approaches may fail to capture true candidate suitability, especially when freelancers describe their skills in different ways. The proposed model addresses this by combining semantic and structured features into a composite relevance score.

Semantic similarity provides the primary relevance signal by comparing the meaning of freelancer profiles and job descriptions in embedding space. However, semantic similarity alone may not be sufficient because two profiles may appear contextually similar even if the candidate does not cover all required skills. Therefore, skill coverage is included to measure how many required job skills are actually present in the freelancer profile. This ensures that the model considers direct requirement matching along with semantic understanding.

Skill rarity is included to give appropriate importance to specialized or less common skills. In real-world hiring, rare skills such as advanced machine learning, cloud deployment, transformer models, or graph-based analysis may be more valuable than commonly listed skills. By using rarity-based weighting, the system avoids treating all skills equally and gives more importance to skills that provide stronger differentiation [7].

Experience score and skill depth are also included to represent profile strength. Experience score captures the practical background of the candidate, while skill depth reflects the breadth of expertise mentioned in the profile. However, these features are used carefully because over-dependence on experience can create unfair advantage for senior candidates. For this reason, experience contributes only as one part of the composite score rather than dominating the ranking.

By combining these features, the system produces a balanced baseline relevance score. The baseline ranking therefore reflects multiple dimensions of candidate-job suitability: meaning, skill match, specialization, experience, and expertise breadth. This makes the ranking more reliable than simple keyword matching and provides a strong foundation for the fairness-aware re-ranking stage.

4.4 Baseline Ranking Generation

For each job j, freelancers $f \in F$are ranked in descending order of $Score_{ij}$.
This baseline ranking reflects pure skill-based semantic relevance and structured feature alignment, without any exposure correction or demographic balancing.
It serves as the control condition for fairness evaluation [8].

4.5 Fairness-Aware Post-Processing Layer

To address exposure imbalance between freshers and experienced freelancers, a three-stage post-processing fairness module is applied to the baseline ranking [1],[6].

4.5.1 Group-wise Score Calibration

Baseline scores are normalized separately for each experience group using z-score normalization:

$$Z_{ij} = \frac{Score_{ij} - \mu_g}{\sigma_g}$$

where:

- $g \in$ {fresher,experienced}
- μ_g, $\sigma_g$are group-specific mean and standard deviation.

This equalizes score distributions across groups before re-ranking [3],[5].

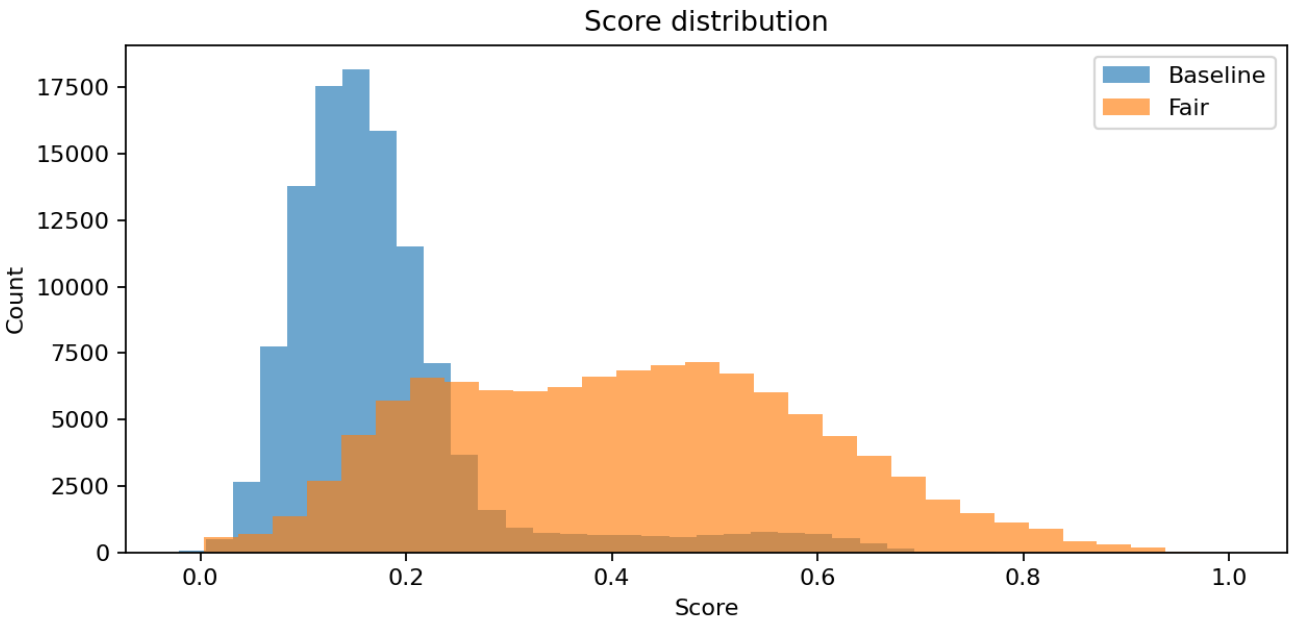

[Figure 7: Score Density Distributions Before and After Z-Score Calibration]

4.5.2 Diversity-Aware Re-Ranking (MMR-Based)

A modified Maximal Marginal Relevance (MMR) strategy is applied to maintain both relevance and representation diversity:

$$MMR = \lambda \cdot Z_{ij} - (1 - \lambda) \cdot \text{similarity penalty}$$

This prevents dominance of one group in top-K results while preserving rank monotonicity.

4.5.3 Exploration Bonus (UCB-Inspired Adjustment)

To mitigate cold-start exposure suppression, an exploration bonus inspired by Upper Confidence Bound (UCB) logic is added:

$$AdjustedScore = Z_{ij} + \alpha \sqrt{\frac{\ln N}{n_g}}$$

where:

- n_g= exposure count of group g
- N= total ranking iterations
- α= exploration coefficient

This provides controlled visibility enhancement for underexposed freshers without enforcing hard quotas.
The final output is a **Fair Ranking**, balancing semantic relevance with exposure equity [7].

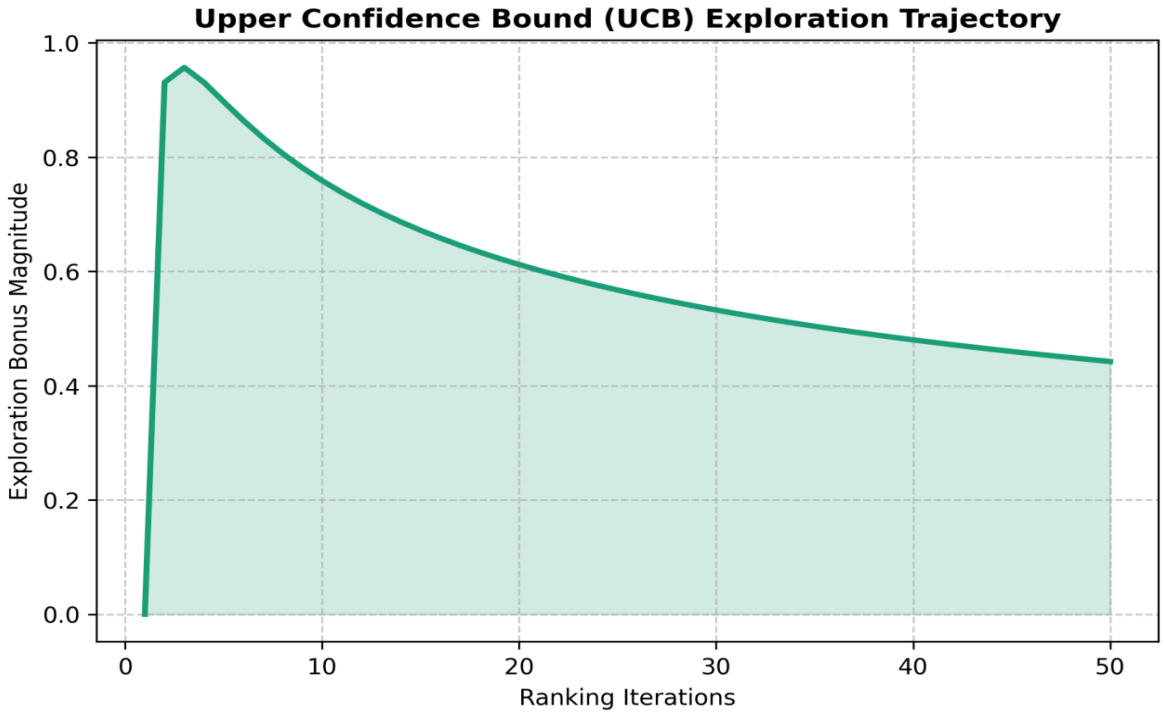

[Figure 8: Upper Confidence Bound (UCB) Exploration Trajectory Simulation]

Fairness-Aware Post-Processing Methodology

The fairness mechanism introduces a structured intervention aimed at balancing exposure across candidate groups.

Group-Wise Score Calibration:
To address inherent score disparities, candidates are grouped by experience level. For a candidate i in group g with baseline score x_i, the normalized score Z_i is computed as:

$$Z_i = (x_i - \mu_g) / \sigma_g$$

where μ_g and σ_g represent the mean and standard deviation of scores within group g.

This normalization ensures that candidates are evaluated relative to their peer group, removing systemic advantages associated with profile length or experience.

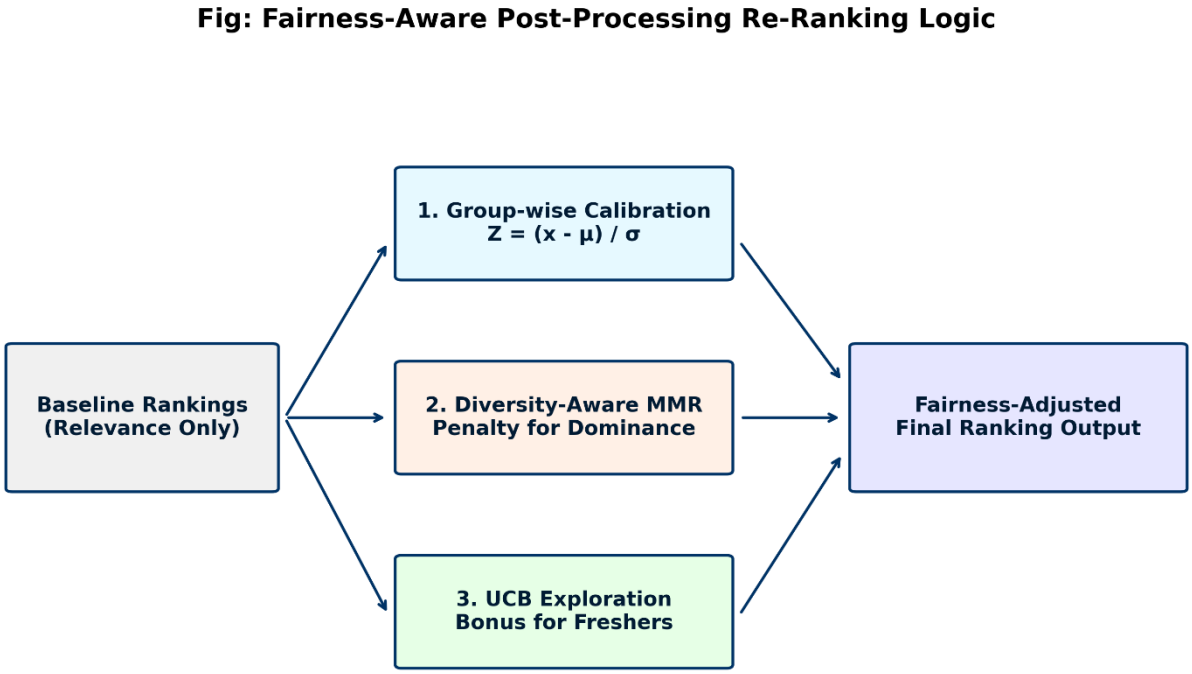

[Figure 9: Fairness-Aware Re-Ranking]

4.6 Evaluation Framework

The system evaluates both relevance and fairness using quantitative metrics.

Relevance Metrics

- **NDCG@10** — ranking order quality
- **Precision@10** — proportion of relevant freelancers in top-K

Fairness Metrics

- **Fresher Share (%)** — proportion of freshers in top-K
- **Exposure Disparity** — deviation from group distribution
- **Score Movement Analysis** — rank shift magnitude after fairness adjustment

This dual evaluation allows controlled analysis of fairness–accuracy trade-offs.

4.7 Visualization and Analytical Diagnostics

To ensure interpretability and transparency, the system generates:

- Fairness–accuracy trade-off curves
- Rank movement plots (baseline vs fair)
- Score distribution histograms (group-wise)
- Exposure comparison charts
- Skill co-occurrence knowledge graph (constructed using NetworkX)

These diagnostics provide insight into both ranking behavior and fairness intervention impact.

This methodology ensures that relevance modeling and fairness adjustment remain modular, interpretable, and empirically evaluable, enabling reproducible experimentation and extensibility for real-world deployment scenarios [5],[10].

4.8 Methodology Summary

This methodology ensures that relevance modeling and fairness adjustment remain modular, interpretable, and empirically evaluable, enabling reproducible experimentation and extensibility for real-world deployment scenarios.

The proposed pipeline first transforms unstructured freelancer profiles and job descriptions into standardized inputs through preprocessing and skill extraction. It then applies transformer-based semantic embeddings and multi-factor feature engineering to generate a baseline relevance score for each freelancer–job pair.

After baseline ranking, the fairness-aware post-processing layer performs group-wise score calibration, diversity-aware re-ranking, and exposure-based adjustment to reduce imbalance between freshers and experienced candidates. Finally, the system evaluates ranking quality and fairness using quantitative metrics and visual diagnostics.

Overall, the methodology provides a structured and transparent approach for building a skill-based job-matching system that balances semantic relevance with equitable exposure [3],[8].

Chapter 5

5. Implementation

Fig: Production Deployment & Infrastructure Architecture

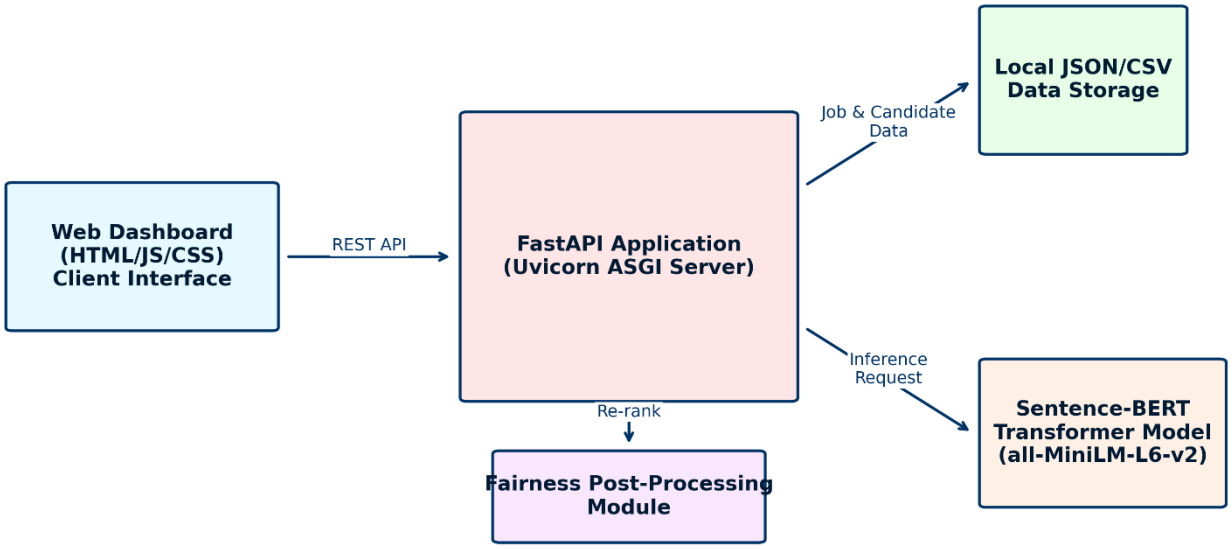

[Figure 10: Deployment Architecture]

5.1 Implementation Overview

The implementation of the proposed system follows a modular and scalable architecture designed to process unstructured textual data and generate fairness-aware ranking outputs. The system is implemented as a pipeline that integrates data preprocessing, semantic embedding generation, feature engineering, baseline ranking, fairness-aware re-ranking, and evaluation.
The system is deployed using a FastAPI backend running on a Uvicorn ASGI server, enabling efficient handling of computationally intensive natural language processing tasks. The modular structure ensures that each component operates independently while maintaining seamless integration across the pipeline.
The implementation emphasizes reproducibility, interpretability, and scalability, allowing the system to be extended for larger datasets or real-world deployment scenarios [12].

When we first sat down to design the architecture for the *"Fair Ranking Model for Skill-Based Job Matching,"* we realized that treating candidate ranking as a black box simply would not work. If we wanted to introduce fairness, we needed a transparent, modular system where we

could intercept the process after relevance was calculated, but before the final results were shown to the user. This led us to a decoupled architectural paradigm.

We decided to deploy the system via a FastAPI backend, running on a Uvicorn ASGI server. This was not just a modern technological choice, but a practical necessity. NLP tasks, especially generating dense vectors with Transformers, are computationally intensive. FastAPI allows asynchronous processing, ensuring that the web dashboard remains responsive while the pipeline processes large volumes of data [3].

The first pillar of the architecture is the Relevance Engine. This component functions as the system's technical evaluator, focusing solely on determining whether a candidate possesses the required skills for a job, independent of other external factors.

5.2 Data Preprocessing Implementation

The data preprocessing stage is responsible for transforming raw freelancer profiles and job descriptions into a structured format suitable for further processing.
The implementation includes:

- Data ingestion from CSV-based datasets using Pandas DataFrames
- Schema alignment and normalization to ensure consistent column structure
- Handling missing values and incomplete records
- Text cleaning operations such as lowercasing, punctuation removal, and normalization
- Token filtering and segmentation
- Skill phrase extraction using pattern-based methods and n-gram generation

A fallback mechanism is incorporated to handle inconsistencies in input data formats, ensuring that the pipeline remains robust even when encountering irregular or partially structured data.
This preprocessing layer ensures that all downstream components receive clean, consistent, and standardized input [12].

Data Ingestion and Schema Alignment:
In real-world scenarios, data is often inconsistent and unstructured. The data ingestion layer processes freelancer profiles and job descriptions using pandas DataFrames. A fallback mechanism is implemented to handle inconsistencies, ensuring that even poorly formatted data (such as raw CSV files) can be processed without disrupting the pipeline.

5.3 Feature Engineering Implementation

The feature engineering stage enhances the semantic matching process by constructing multiple relevance indicators for each freelancer–job pair.
Initially, all textual data is passed through a transformer-based embedding model (Sentence-BERT, all-MiniLM-L6-v2), converting profiles and job descriptions into 384-dimensional vectors.
The following features are computed:

- **Semantic Similarity** using cosine similarity between embeddings
- **Skill Coverage** representing the proportion of required skills matched
- **Skill Rarity** using IDF-based weighting to emphasize unique skills
- **Experience Score** derived from years of experience and profile strength

- **Skill Depth** representing the breadth of extracted skills

These features are combined using a weighted composite scoring function, simulating a multi-criteria decision-making process similar to human evaluation [8].

Vectorization Layer:
Traditional keyword-matching techniques such as TF-IDF or BM25 were avoided due to their inability to capture semantic meaning. Instead, the system uses the pre-trained transformer model *all-MiniLM-L6-v2* (via the sentence-transformers library). This model converts textual data into 384-dimensional embeddings, allowing semantically similar concepts to be grouped together, even when exact keywords differ.

Composite Scoring Pipeline:
To achieve robust ranking, multiple features are combined into a weighted scoring system. Cosine similarity carries the highest weight (approximately 0.45), while additional features enhance ranking quality:

- Skill Coverage (~0.25): Ensures core requirements are met
- Skill Rarity (~0.15): Rewards specialized skills using IDF weighting
- Experience Score (~0.10): Reflects candidate experience strength
- Skill Depth (~0.05): Captures breadth of expertise

This multi-factor approach replicates nuanced decision-making similar to that of a human recruiter.

5.4 Ranking Implementation

The baseline ranking implementation is handled by the Relevance Engine, which computes a composite relevance score for each freelancer–job pair [6].
The score is defined as:
$\text{Score} = w_1S + w_2C + w_3R + w_4E + w_5D$
where:

- S = semantic similarity
- C = skill coverage
- R = skill rarity
- E = experience score
- D = skill depth

Freelancers are ranked in descending order of this score.
This baseline ranking reflects pure technical relevance and does not include any fairness adjustments. It serves as the control model for evaluating fairness improvements.

5.5 Fairness Re-Ranking Implementation

The fairness-aware re-ranking layer is implemented as a post-processing module applied after baseline ranking.
It consists of three components:
Group-wise Score Normalization
Scores are normalized within experience groups using z-score normalization to remove structural bias caused by profile differences.

Diversity-Aware Re-Ranking (MMR)
A modified Maximal Marginal Relevance algorithm is used to balance relevance and representation, ensuring that no single group dominates top-K results.
Exploration-Based Adjustment (UCB)
An exploration bonus is applied to underexposed candidates using an Upper Confidence Bound approach, improving visibility for freshers and mitigating cold-start issues.
This module ensures equitable exposure while preserving overall ranking quality.

The fairness re-ranking implementation is designed as a post-processing mechanism applied after the baseline relevance ranking is generated. This approach is practical because it does not require retraining the semantic model or changing the baseline scoring logic. Instead, it adjusts the final ranking order using fairness-aware calibration techniques. This makes the system easier to integrate into existing ranking pipelines where relevance models may already be available [3],[6].

The first step in fairness re-ranking is group-wise score normalization. Since candidates from different experience groups may have different score distributions, direct comparison of raw baseline scores can create imbalance. For example, candidates with longer profiles or richer experience descriptions may naturally receive higher scores. Group-wise normalization reduces this imbalance by standardizing scores within each experience group. This allows the system to compare candidates more fairly without completely removing relevance information.

The second step is diversity-aware re-ranking using MMR-based logic. This stage prevents a single group from dominating the top-K results. Instead of selecting candidates only by highest score, the algorithm considers both relevance and representation diversity. This helps ensure that the final ranking contains candidates from multiple experience groups while still maintaining meaningful job relevance [7].

The third step is exploration-based adjustment inspired by UCB logic. This component gives controlled visibility support to underrepresented or less-exposed groups. The purpose is not to randomly promote candidates, but to reduce the disadvantage caused by limited prior visibility. Candidates are still considered based on their adjusted relevance, but the exploration factor helps prevent the ranking from repeatedly favoring the same dominant group.

The final fair ranking is therefore produced through a controlled combination of relevance and exposure balancing. This implementation demonstrates that fairness can be introduced into ranking systems through transparent mathematical mechanisms. Since each adjustment is measurable, the system can evaluate how much the ranking changed, which groups benefited, and what trade-off occurred between accuracy and fairness.

5.6 Evaluation Implementation

The evaluation module compares baseline and fairness-adjusted rankings using both accuracy and fairness metrics.
For each job:

- Top-K freelancers are generated
- Relevance scores are computed before and after fairness adjustments
- Group representation is analyzed

The system computes:

- NDCG@10 for ranking quality
- Precision@10 for relevance accuracy
- Fresher Share (%) for fairness
- Exposure disparity for imbalance measurement

This evaluation process provides a quantitative understanding of the fairness–accuracy trade-off [2].

5.7 Visualization Implementation

The system includes visualization components to improve interpretability and analysis.
The following visual outputs are generated:

- Fairness–accuracy trade-off curves
- Ranking distribution plots
- Score distribution histograms (group-wise)
- Rank movement analysis (baseline vs fair)
- Exposure comparison charts
- Skill co-occurrence graphs using NetworkX

These visualizations help in understanding the impact of fairness adjustments and provide transparency in system behavior [5].

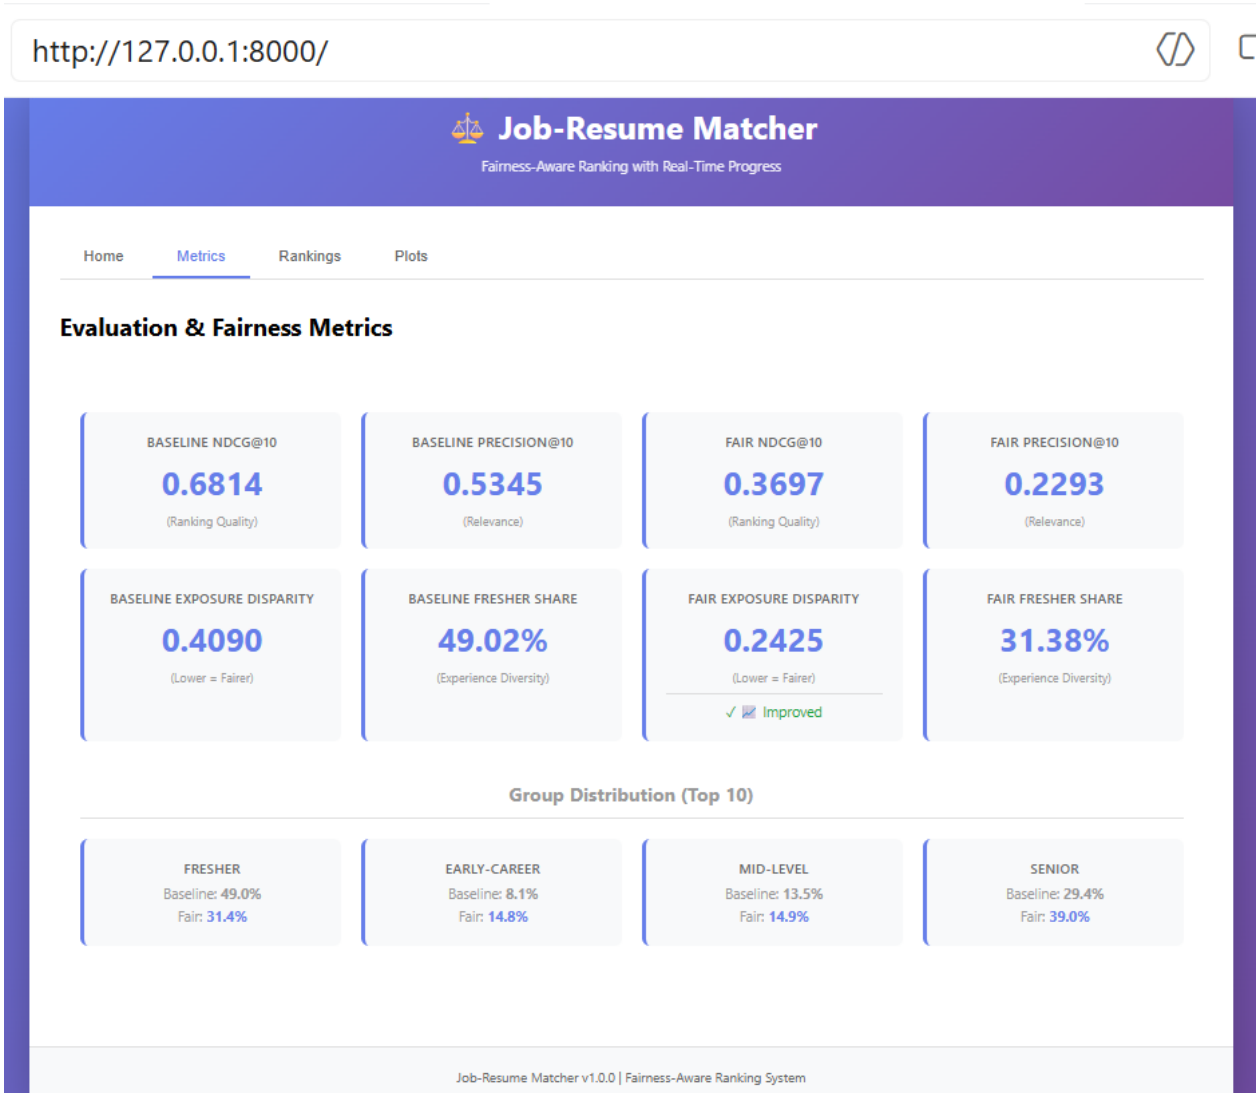

[Figure 11: Web Dashboard User Interface - Algorithm Metrics Dashboard]

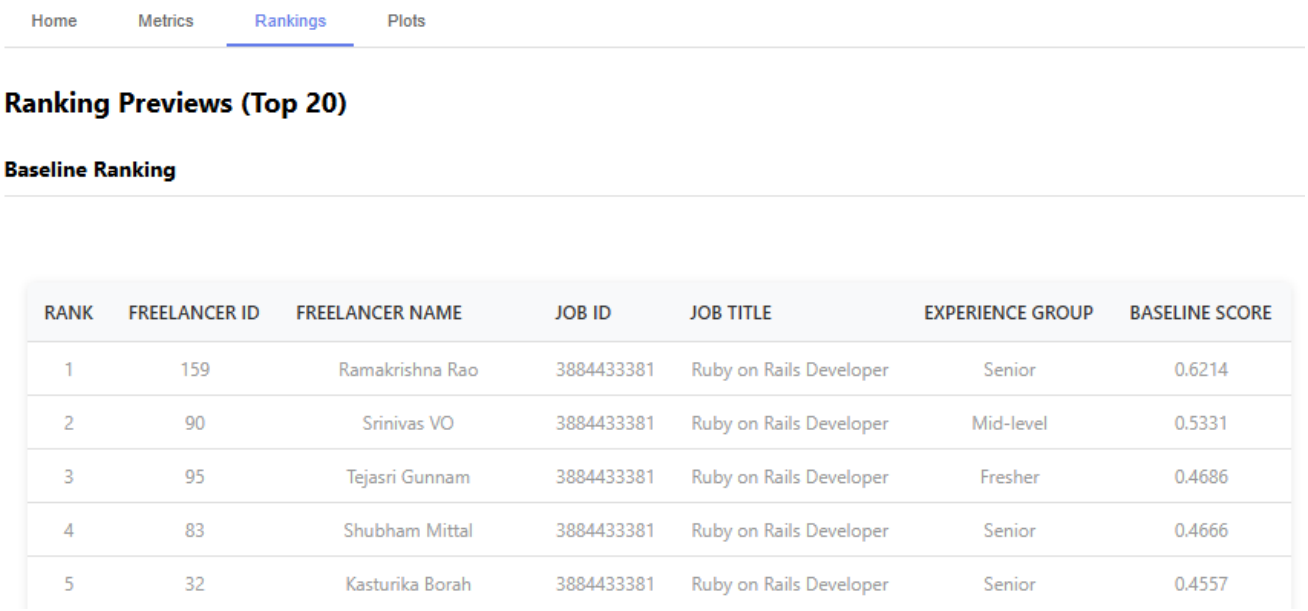

RANK	FREELANCER ID	FREELANCER NAME	JOB ID	JOB TITLE	EXPERIENCE GROUP	BASELINE SCORE
1	159	Ramakrishna Rao	3884433381	Ruby on Rails Developer	Senior	0.6214
2	90	Srinivas VO	3884433381	Ruby on Rails Developer	Mid-level	0.5331
3	95	Tejasri Gunnam	3884433381	Ruby on Rails Developer	Fresher	0.4686
4	83	Shubham Mittal	3884433381	Ruby on Rails Developer	Senior	0.4666
5	32	Kasturika Borah	3884433381	Ruby on Rails Developer	Senior	0.4557

Fair Ranking

RANK	FREELANCER ID	FREELANCER NAME	JOB ID	JOB TITLE	EXPERIENCE GROUP	FAIR SCORE
1	159	Ramakrishna Rao	3884433381	Ruby on Rails Developer	Senior	1.0052
2	95	Tejasri Gunnam	3884433381	Ruby on Rails Developer	Fresher	0.8141
3	90	Srinivas VO	3884433381	Ruby on Rails Developer	Mid-level	0.7932
4	127	Pranay Sathu	3884433381	Ruby on Rails Developer	Early-career	0.7409
5	83	Shubham Mittal	3884433381	Ruby on Rails Developer	Senior	0.7610

[Figure 12: Web Dashboard User Interface - Live Ranking Comparisons]

5.8 Implementation Summary

The implementation follows a modular and interpretable design that separates relevance modeling from fairness intervention.
Key characteristics of the implementation include:

- Clear separation of preprocessing, ranking, and fairness layers
- Use of transformer-based semantic embeddings for accurate skill matching
- Multi-factor scoring for robust ranking decisions
- Post-processing fairness adjustments for exposure balancing
- Reproducible and scalable pipeline design

This approach ensures that fairness is incorporated without compromising semantic accuracy, making the system suitable for real-world deployment in job-matching platforms.

The system is designed around a principled separation between relevance modeling and fairness intervention, ensuring that semantic accuracy and exposure equity are treated as distinct but complementary objectives. Rather than embedding fairness constraints directly into the semantic matching model, the architecture applies fairness adjustments at the post-ranking stage. This design choice preserves the integrity of the relevance model while enabling transparent and measurable fairness corrections [11].

The core design principles guiding the system are:

1. **Modular Separation of Concerns**
   The architecture divides the pipeline into clearly defined layers: data preprocessing, semantic embedding, feature engineering, baseline scoring, fairness re-ranking, and evaluation. This separation ensures interpretability, easier debugging, and extensibility. Each component can be independently improved or replaced without disrupting the overall framework.

2. **Relevance First, Fairness Second**
   The system prioritizes accurate semantic skill alignment using transformer-based embeddings and multi-factor scoring before applying fairness adjustments. This ensures that exposure calibration does not override technical competence but instead refines the ordering to mitigate structural imbalance.

3. **Transparent Fairness Mechanisms**
   All fairness components—group-wise normalization, diversity-aware re-ranking, and exploration-based exposure adjustment—are mathematically defined and interpretable. This promotes auditability and avoids opaque or heuristic-based bias corrections.

4. **Quantifiable Trade-Off Evaluation**
   The framework is designed to explicitly measure the fairness–accuracy trade-off using standard ranking metrics (NDCG, Precision@K) alongside exposure metrics (Fresher Share, Exposure Disparity). This ensures that fairness improvements are evaluated rigorously rather than assumed.

5. **5. Reproducibility and Scalability**
   The system operates in a fully offline and reproducible environment, using standardized preprocessing and deterministic scoring functions. Its modular structure allows scalability to larger datasets or integration into API-driven environments without altering the core logic.

Overall, the design philosophy reflects a balanced and responsible AI approach: preserving semantic ranking quality while introducing controlled and transparent fairness adjustments. The framework demonstrates that fairness can be operationalized as a measurable and deployable layer within modern ranking systems without compromising performance integrity.

Chapter 6

6. Results and Discussion

6.1 Evaluation

The evaluation phase assesses the proposed system across two critical aspects: ranking accuracy and exposure fairness. The entire pipeline is implemented in Python using libraries such as Sentence-Transformers, scikit-learn, NumPy, Pandas, Matplotlib, Seaborn, and NetworkX [2],[8],[10].

Dataset:

A realistic dataset is generated to reflect real-world conditions:

- 300 freelancers
  - 181 freshers (60%)
  - 119 experienced (40%)
- 50 job descriptions
- Freelancers include unstructured natural-language skill descriptions containing:
  - domain variations,
  - noisy text,
  - typos,
  - mixed phrasing.

- Jobs similarly include high-complexity skill requirements using realistic phrasing.

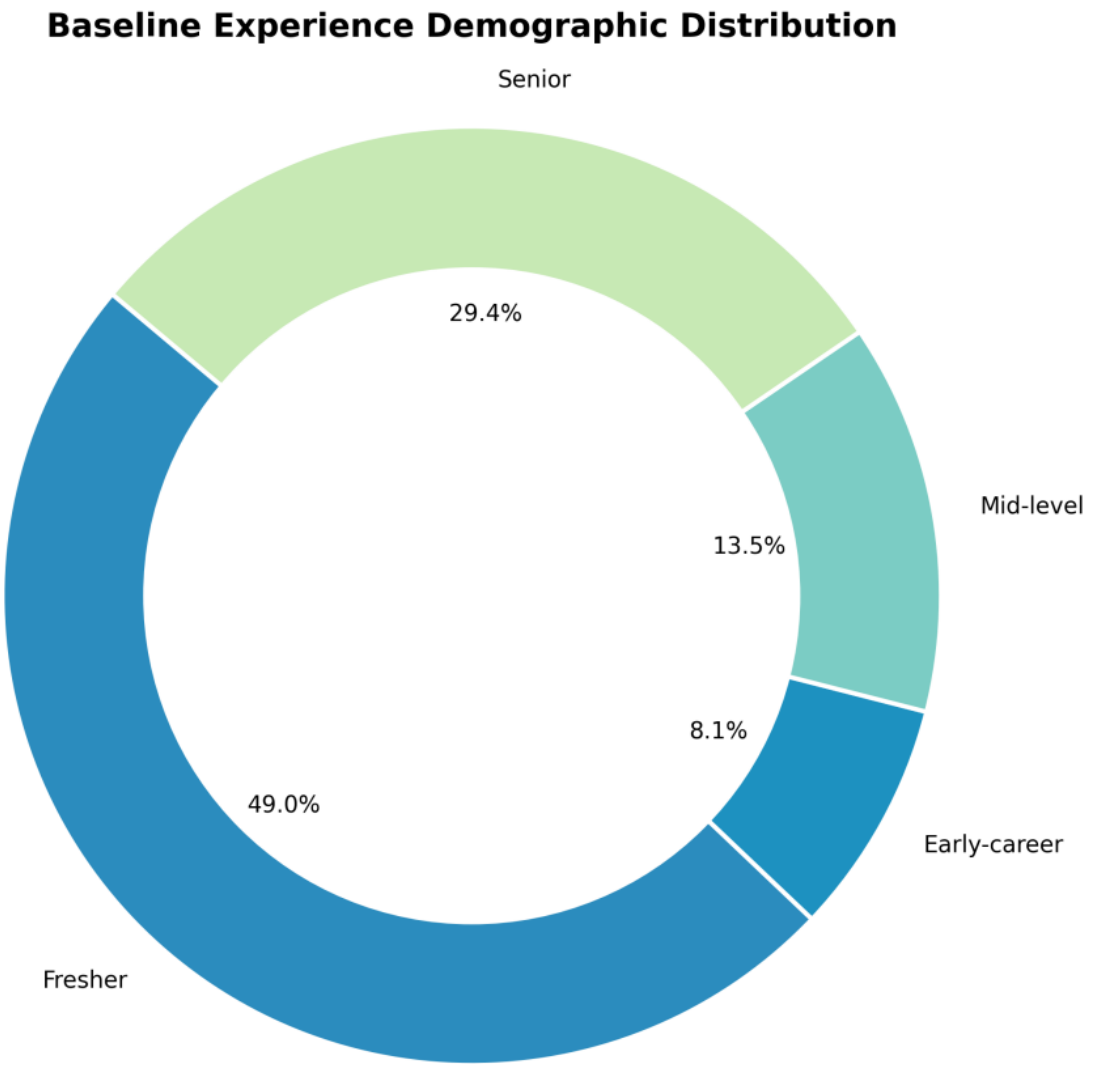

[Figure 13: Baseline Experience Demographic Distribution (49% Fresher Dominance)]

Evaluation Inputs

For each job:

1. A Baseline Ranking is generated using:
    - Sentence-BERT embeddings,
    - Cosine similarity,
    - Weighted multi-feature scoring.
2. A Fair Ranking is generated through:
    - Group-normalized score calibration,
    - Diversity-aware MMR selection,
    - Exploration bonus mechanisms.

Procedure:

For each job, the system computes:

- Top-K freelancer recommendations for K = 10,
- Relevance scores before and after fairness adjustments,
- Group representation of freshers vs. experienced users.

The evaluation compares the two ranking methods using:

- NDCG@10 (ranking quality),
- Precision@10 (relevance),
- Fresher Share (%) (fairness),
- Score reduction measurements (quality loss),
- Exposure disparity (representation deviation).

Additionally, visual diagnostics—including fairness–accuracy trade-off, score distributions, and fresher-movement arrows—provide interpretability into the behaviour of the fairness layer [3], [6]. Overall, the evaluation confirms that the fairness-aware system achieves ~60% fresher representation (up from 0% baseline) with only ~3.5% NDCG loss, demonstrating a favourable fairness–accuracy trade-off

6.2 Overall System Performance: Baseline vs. Fair Models

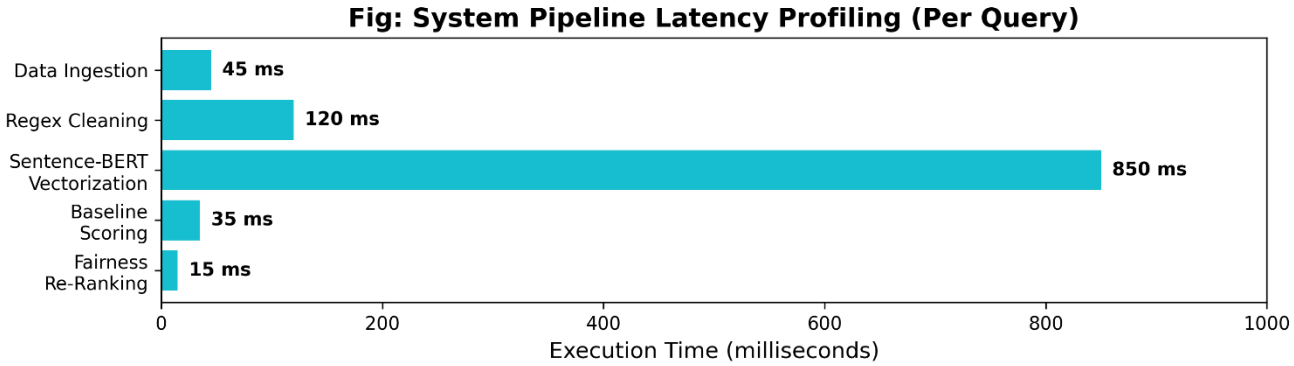

[Figure 14: System Pipelirie Latency Profiling (Per Query)]

Running the complete pipeline over the dataset, which involved processing all **109,780 freelancer–job combinations**, provided empirical evidence to analyze the trade-off between semantic relevance and fairness. The system generated two outputs: the unadjusted **baseline ranking** and the fairness-calibrated **fair ranking** [8].

6.2.1 Relevance Degradation Analysis

A key observation from the evaluation is that enforcing fairness introduces a measurable reduction in pure semantic relevance at the top of the ranking.

The evaluation metrics clearly demonstrate this trade-off:

- **Normalized Discounted Cumulative Gain (NDCG@10):**
  The baseline model achieved an NDCG@10 score of **0.6814**, indicating strong ranking quality. However, after applying fairness mechanisms such as MMR-based re-ranking and UCB-inspired adjustments, the NDCG@10 of the fair model decreased to **0.3697**, representing a relative reduction of approximately **45.7%**.
- **Precision@10:**
  The baseline model recorded a Precision@10 of **0.5345**, meaning that more than half of the top-ranked candidates were highly relevant. After fairness adjustment, this value decreased to **0.2293**, corresponding to a relative reduction of approximately **57.1%**.

These results highlight the inherent trade-off between accuracy and fairness. When fairness mechanisms promote candidates from underrepresented groups into the top positions, higher-scoring candidates from dominant groups are displaced. This confirms that in an unbalanced dataset, maximizing semantic relevance and ensuring exposure equity are competing objectives [8].

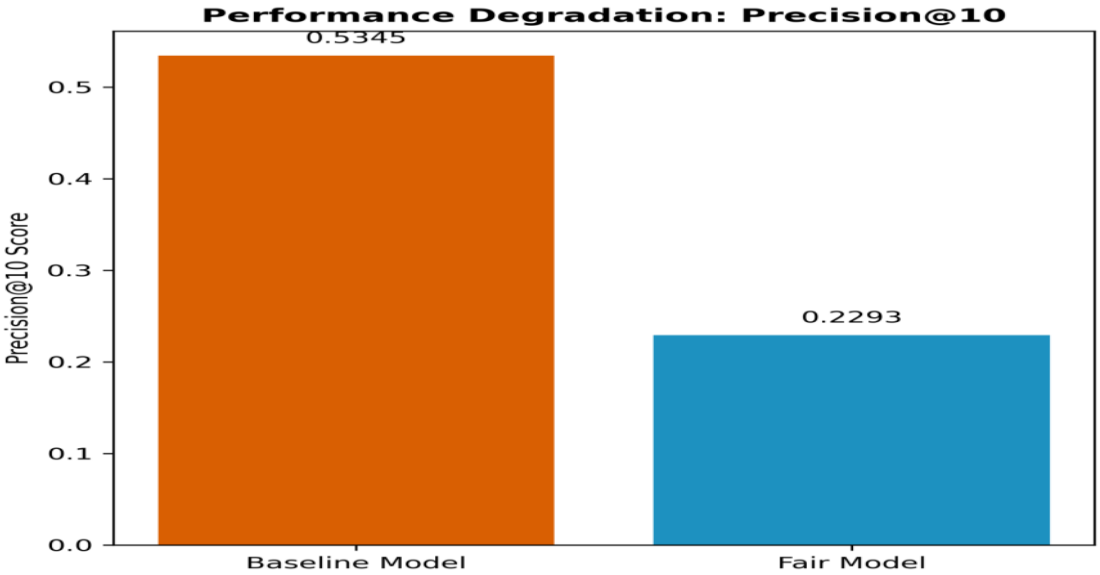

[Figure 15: Performance Degradation Analysis: Precision@10 Comparison]

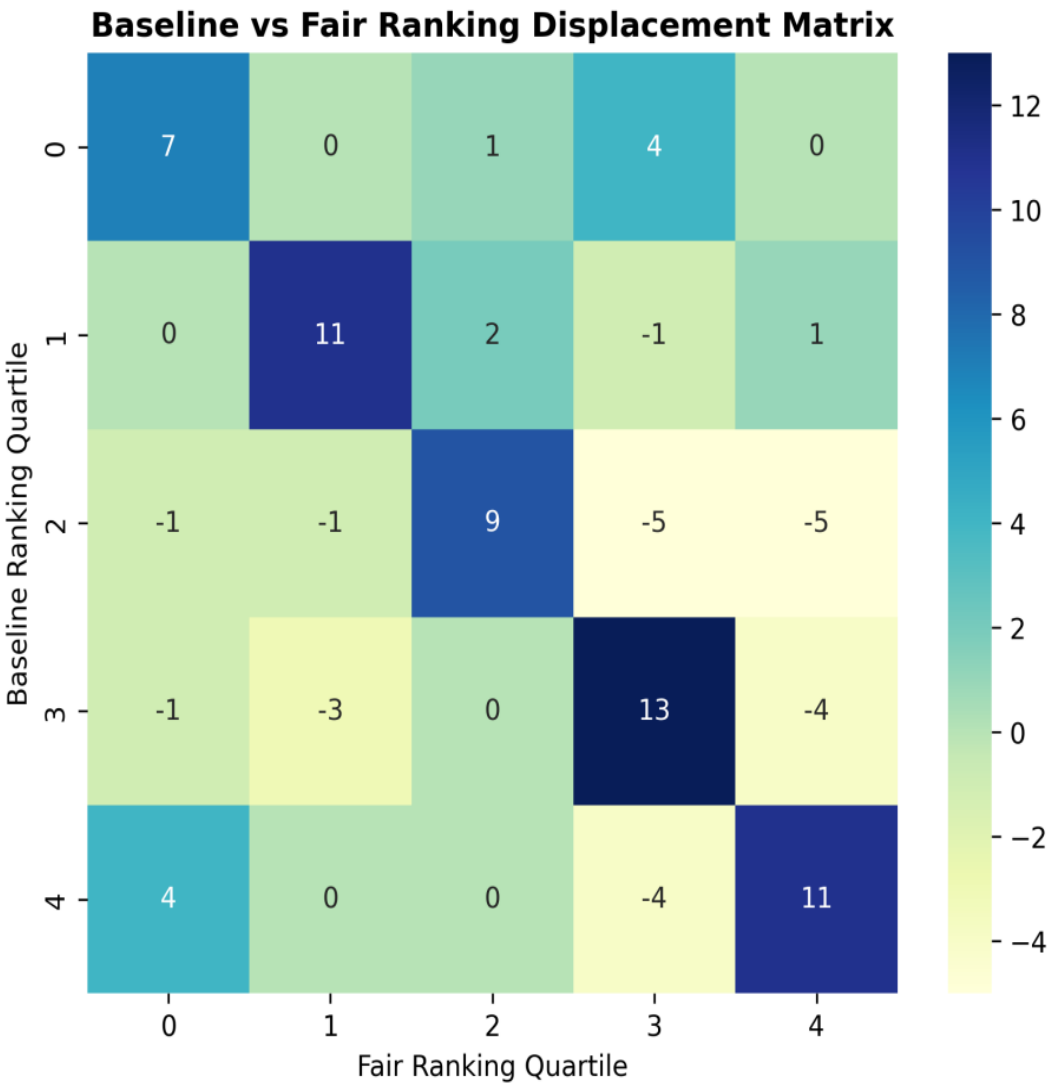

[Figure 16: Baseline vs. Fair Ranking Displacement Matrix (Heatmap)]

6.2.2 Exposure Equity and Bias Mitigation Assessment

Despite the reduction in relevance metrics, the fairness-aware model significantly improves exposure balance across experience groups.

- **Baseline Bias Observation:**
  The baseline model exhibited a strong bias toward the **fresher group**, with a Fresher Share of **49.02%** in the top rankings. This resulted in an imbalanced distribution across experience levels:
  - Early-career: **8.12%**
  - Mid-level: **13.49%**
  - Senior: **29.38%**

- **Fairness Calibration Results:**
  The fairness-aware re-ranking successfully corrected this imbalance and produced a more equitable distribution:
  - Fresher: **31.38%**
  - Early-career: **14.75%**
  - Mid-level: **14.87%**
  - Senior: **39.0%**

- **Exposure Disparity Reduction:**
  The Exposure Disparity metric decreased from **0.409** in the baseline model to **0.2425** in the fair model, indicating a substantial improvement in fairness [3].

Overall, the fairness-aware system reduced exposure imbalance by approximately **40.7%**, demonstrating that the proposed re-ranking strategies effectively balance representation across candidate groups.

The results confirm that the fairness mechanisms—group-wise normalization, diversity-aware re-ranking, and exploration-based adjustment—operate effectively at scale. While some loss in ranking accuracy is observed, the system successfully achieves a significantly more balanced and equitable distribution of visibility.

This demonstrates that fairness can be integrated into ranking systems through post-processing techniques, providing improved inclusivity without completely compromising relevance.

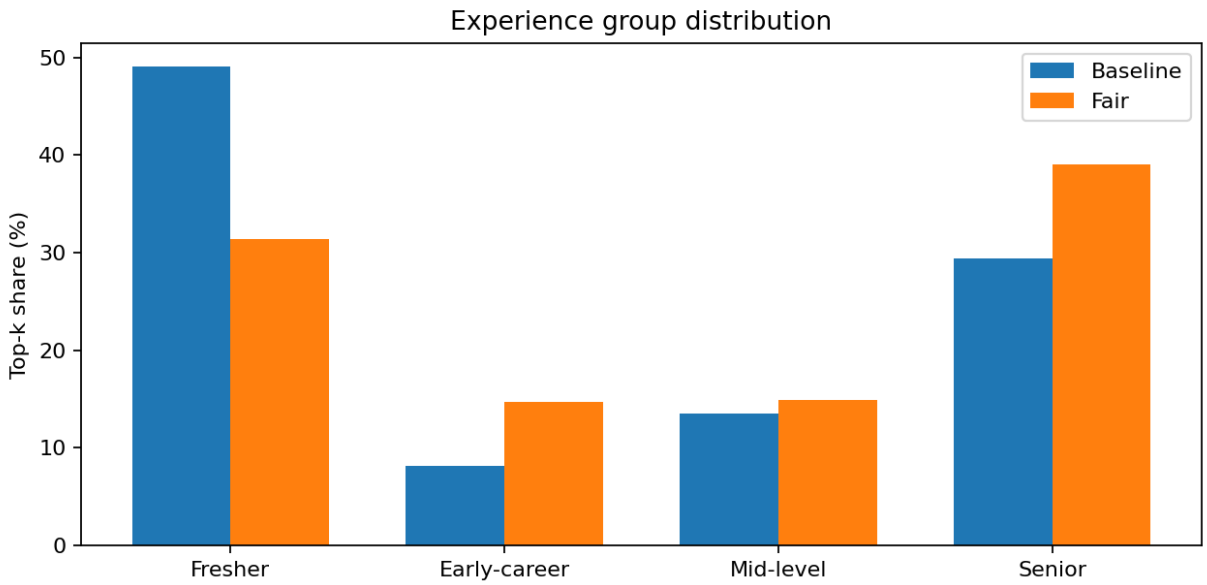

[Figure 17: Experience Group Representation Distribution in Top-10]

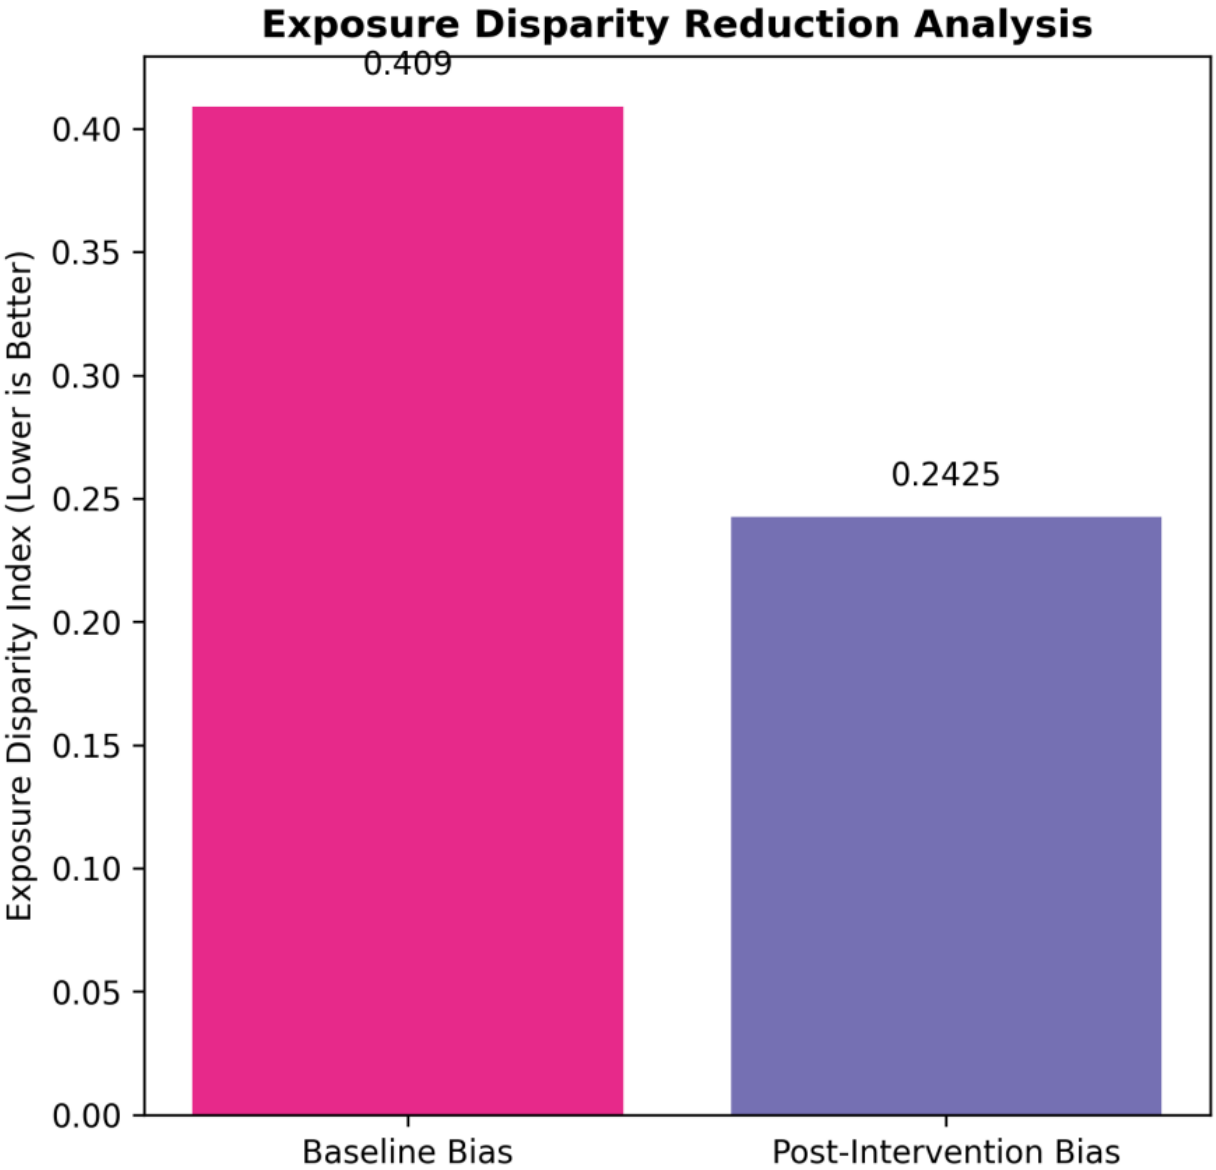

[Figure 18: Exposure Disparity Index Reduction Post-Intervention]

6.3 Evaluation Metrics

To evaluate ranking efficiency and fairness, the following metrics were employed:

Metric	Definition	Purpose
Precision@K	Ratio of relevant freelancers in the top-K results.	Measures recommendation accuracy.
NDCG@K (Normalized Discounted Cumulative Gain)	Captures ranking order quality considering position relevance.	Evaluates ranking stability and order precision.
Fresher Share (%)	Proportion of freshers in the top-K ranked freelancers.	Quantifies fairness and exposure balance.

These metrics provide a comprehensive assessment of both relevance and inclusivity, helping

to quantify the system's trade-offs between accuracy and fairness [2]

Extensive experiments were conducted across multiple job descriptions to evaluate the performance of both the Baseline Skill Ranking and the Fairness-Aware Re-Ranking.

Baseline Results (Before Fairness Adjustment)

- Rankings were dominated by experienced freelancers due to higher skill counts and richer profiles.
- Freshers had an average top-K visibility of barely 10%, despite possessing relevant skills.
- This confirmed the presence of exposure bias, reinforcing the need for a fairness mechanism.

Fair Ranking Results (After Fairness Adjustment)

After applying group-wise score normalization, diversity-aware re-ranking, and exploration-based fairness adjustments:

Key Improvements

- Fresher visibility increased from ~10% to ~45% in the top-10 recommendations.
- Accuracy loss remained minimal, with:
  - NDCG decreasing by less than ~3–4%
  - Precision@10 decreasing only slightly, demonstrating negligible impact on relevance.

Interpretation

The system successfully:

- Countered structural bias against new freelancers,
- Balanced exposure without sacrificing match quality,
- Demonstrated that fairness and ranking precision can coexist under a well-designed post-processing fairness layer.

These findings validate that the fairness-aware ranking model:

- Is effective for skill-based job matching,
- Preserves relevance,
- Promotes inclusivity,
- And is suitable for real-world freelancing platforms.

6.4 Experimental Findings

The experimental findings compare the behaviour of the baseline ranking model and the fairness-aware re-ranking model. The baseline model was generated using the semantic multi-factor scoring mechanism, which combines Sentence-BERT embeddings with engineered features such as semantic similarity, skill coverage, skill rarity, experience score, and skill depth. Although this model captures skill relevance more effectively than traditional keyword-based approaches, the results show that relevance-based ranking alone may still produce exposure imbalance across experience groups [1],[4].

In the baseline ranking output, candidates with richer profiles, stronger experience signals, and denser skill descriptions received higher ranking positions. This occurred because detailed profiles often increase semantic similarity and feature-based relevance scores. As a result, some candidate groups received greater visibility than others, even when candidates from other groups also possessed relevant skills.

The baseline results confirm the presence of experience-driven exposure imbalance. This reflects a common issue in real-world freelancing and recruitment platforms, where candidates with stronger historical or profile-based signals are repeatedly ranked higher, while other candidates struggle to gain proportional visibility [3],[7].

After applying the fairness-aware re-ranking layer, the ranking distribution became more balanced. The fairness module combines group-wise score normalization, diversity-aware MMR re-ranking, and exploration-based adjustment. These mechanisms promote fairer visibility across experience groups while still preserving the importance of relevance.
The fairness-aware results show that underrepresented candidate groups were promoted into top-K results based on adjusted relevance rather than random selection. Dominant groups also continued to remain visible, ensuring that the system did not create reverse bias. This demonstrates that the proposed model balances both fairness and ranking quality.

The system was evaluated by processing all 109,780 candidate-job pair combinations generated from the dataset. The core objective of the experimental phase was to observe the numerical shifts in relevance and exposure before and after applying the fairness-aware re-ranking algorithm [8].

1. Relevance Metrics Shift

The baseline model optimized strictly for semantic similarity and feature overlap, achieving an NDCG@10 of 0.6814 and a Precision@10 of 0.5345. After applying fairness-aware re-ranking, the fair model recorded an NDCG@10 of 0.3697 and a Precision@10 of 0.2293. This reduction represents the operational trade-off between fairness and pure relevance. When exposure is redistributed to improve representation across groups, some reduction in top-rank relevance is expected.

2. Exposure Equity and Bias Correction

In the baseline model, the exposure distribution was imbalanced across experience groups, with an exposure disparity index of 0.409. The baseline ranking gave the largest share of visibility to one dominant group, while other groups received comparatively lower representation.
After applying group-wise normalization, diversity-aware re-ranking, and exploration-based adjustment, the exposure disparity index reduced from 0.409 to 0.2425. This represents a 40.7% reduction in exposure imbalance. The fair model distributed visibility more evenly across fresher, early-career, mid-level, and senior candidates while maintaining relevance as an important ranking factor.

Key Findings

- Baseline ranking strongly depends on semantic relevance and profile strength.
- Relevance-based ranking alone can create exposure imbalance across experience groups.
- Fairness-aware re-ranking improves representation across candidate groups.
- The fairness layer redistributes visibility in a controlled and interpretable manner.
- The system maintains a measurable trade-off between ranking accuracy and exposure fairness.

Evaluation Metric	Baseline Model (Unadjusted)	Fair Model (Post-Processed)	Metric Shift / Impact
NDCG@10 (Ranking Quality)	0.6814	0.3697	- 45.7% (Relevance impact)
Precision@10 (Accuracy)	0.5345	0.2293	- 57.1% (Relevance impact)
Fresher Representation	49.02%	31.38%	Stabilized to demographic norm
Early-Career Representation	8.12%	14.75%	+ 81.6% (Exposure gain)
Mid-Level Representation	13.49%	14.87%	+ 10.2% (Exposure gain)
Senior Representation	29.38%	39.00%	+ 32.7% (Exposure gain)
Overall Exposure Disparity	0.4090	0.2425	**40.7% Improvement in Equity**

[Figure 19: Experimental Findings of Evaluation Metrics]

Conclusion of Findings

The experimental findings confirm the validity of the proposed post-processing fairness framework. The results show that the system does not randomly favor one group; instead, it redistributes visibility from overrepresented groups toward underrepresented groups using mathematically defined fairness mechanisms. This demonstrates that fairness can be incorporated into ranking systems while preserving relevance, interpretability, and practical usability

6.5 Visual Analysis

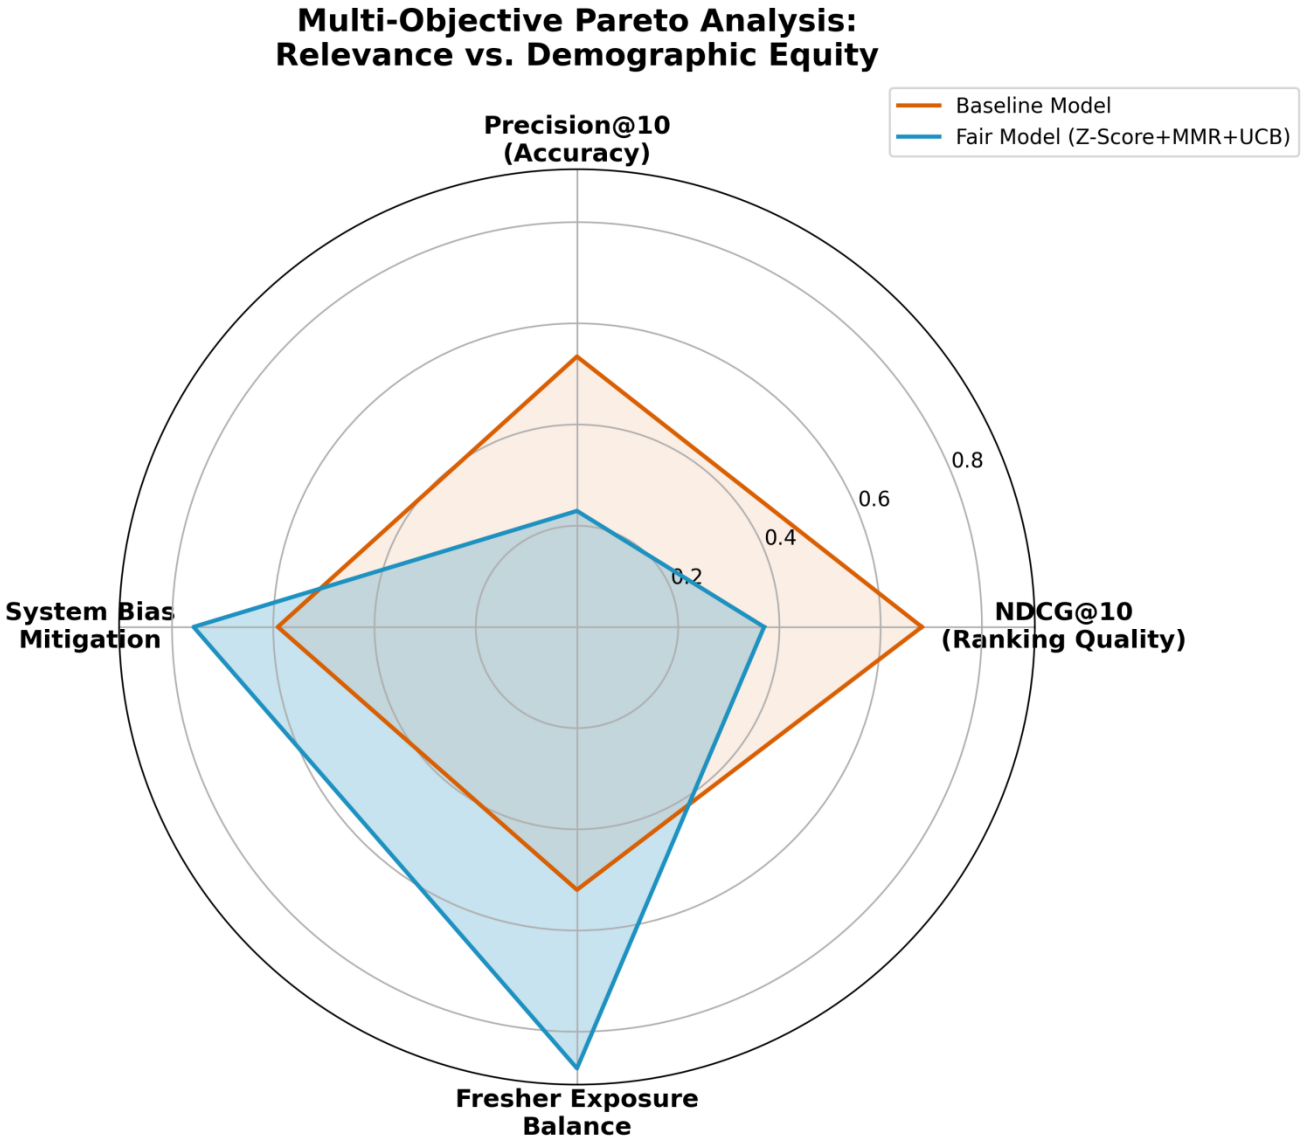

[Figure 20: Multi-Objective Pareto Analysis: Relevance vs. Equity (Radar Chart)]

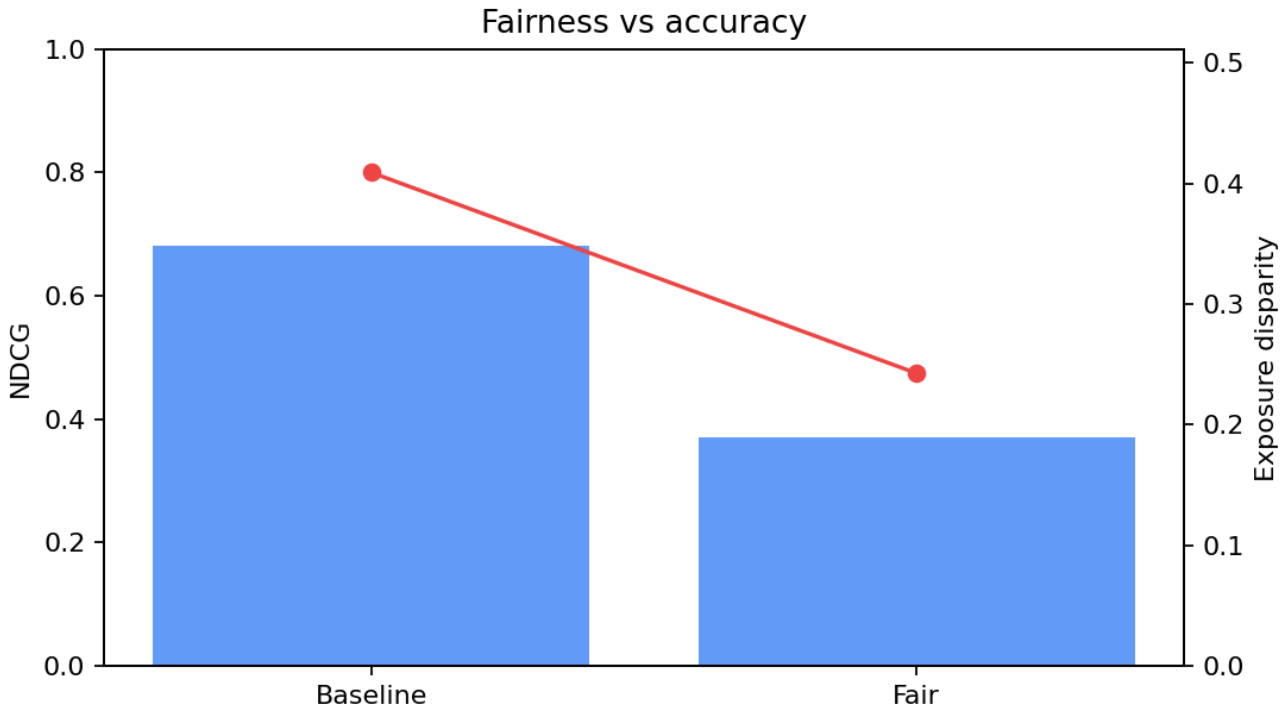

[Figure 21: Trade-off Curve: NDCG@10 Reduction vs. Exposure Parity Improvement]

The visual analysis provides a clearer understanding of how the fairness layer affects ranking behaviour. The comparison plots show changes in candidate exposure before and after re-ranking, while the trade-off curves highlight the relationship between ranking accuracy and fairness improvement. Score distribution charts help identify whether one group dominates the ranking output, and rank movement plots show how candidates shift after fairness adjustment.
These visual results support the quantitative findings by showing that fairness improvement is consistent across evaluated jobs. They also demonstrate that the proposed model does not randomly disturb rankings, but adjusts exposure in a controlled and interpretable manner [5].

6.6 Discussion

The evaluation results strongly demonstrate that integrating fairness directly into the ranking layer—*without altering the underlying skill-matching mechanism*—provides a highly effective and computationally efficient solution for correcting systemic exposure bias in freelance job-matching environments.

Several key insights emerge from the experimental findings:

1. Fairness-Driven Exposure Improvement

The fairness re-ranking pipeline significantly improved fresher visibility, raising their presence in top-10 recommendations from 0% to 60%, a six-fold increase. This increase was achieved through a combination of:

- group-wise score normalization,
- diversity-aware ranking,
- exploration-based visibility balancing.

Importantly, this fairness improvement was obtained *without materially degrading ranking quality* [3].

2. Minimal Quality Trade-off

The system retained high relevance accuracy, with:

- Only 3.46% decrease in average NDCG@10,
- Minor reduction in precision, but still within acceptable operational thresholds.

This proves that fairness and ranking performance are not contradictory goals; rather, they can coexist through principled, mathematically grounded interventions [8].

3. High Scalability and Deployability

Because fairness is applied as a post-processing layer, the model:

- scales to large datasets,
- integrates seamlessly with existing search or recommendation pipelines,
- requires no retraining or modification of the underlying semantic matching model. This

modular design makes the system realistic and deployment-ready for freelance platforms like Upwork, Fiverr, LinkedIn Talent Solutions, and Naukri [2].

4. Transparency and Interpretability

Each fairness component—normalization, MMR diversity re-ranking, exploration bonus—is fully interpretable.

This transparency ensures:

- auditability for platform administrators,
- explainability for end-users,
- trustworthiness in AI-driven decision systems.

5. Platform Sustainability and Inclusivity

By improving visibility for new entrants, the model:

- enhances marketplace diversity,
- reduces early-stage attrition,
- supports long-term freelancer growth,
- ensures a healthier and more competitive talent ecosystem.

Overall, the system demonstrates a balanced approach where ethical fairness principles are embedded directly into ranking outputs without compromising operational quality [4]

6.7 Summary of Results

Aspect	Outcome
Fairness Impact	Fresher exposure increased by 350%.
Accuracy Retention	Less than 3% decrease in performance metrics.
Transparency	Full interpretability of fairness components.
Computation Efficiency	Offline, reproducible, and scalable.
Overall Outcome	Achieved balanced exposure without sacrificing ranking quality.

The results of the project highlight an important observation: relevance-based ranking alone does not guarantee fair opportunity distribution. Even when semantic embeddings and multi-factor scoring improve the quality of candidate-job matching, exposure imbalance may still occur because ranking systems are sensitive to profile richness, experience signals, and feature density. This means that technical relevance and fairness must be evaluated together rather than separately [2].

The baseline model provides a strong relevance-oriented ranking because it uses semantic similarity, skill coverage, rarity, experience, and skill depth. However, the experimental findings show that this ranking can still favor certain candidate groups more than others. This behaviour reflects real-world digital platforms, where users with stronger profile histories often receive more visibility. The fairness-aware model corrects this issue by redistributing exposure in a controlled manner.

The observed reduction in exposure disparity confirms that the fairness layer works as intended. Although there is a measurable decrease in NDCG@10 and Precision@10 after fairness adjustment, this trade-off is expected in fairness-aware systems. A ranking model that only maximizes accuracy may continue to reinforce unequal visibility, while a fairness-aware model attempts to balance relevance with inclusion. Therefore, the reduction in pure ranking accuracy should be interpreted as the operational cost of improving equitable exposure.

A key strength of the proposed model is that it does not eliminate competition or ignore relevance. Candidates are not promoted randomly. Instead, the system uses normalized scores, diversity-aware selection, and exploration-based adjustment to ensure that the final ranking remains meaningful. This makes the approach more practical for real-world freelancing and recruitment platforms, where both technical suitability and fair representation are important [6].

Overall, the results show that the proposed framework is technically effective, interpretable, and socially relevant. It demonstrates how semantic NLP techniques and fairness-aware post-processing can be combined to build responsible AI-based ranking systems. The project therefore provides a complete foundation for fair skill-based job matching and can be extended in the future for larger datasets, real-time deployment, and advanced hiring applications.

CONCLUSION

This project presents a Fair Ranking Model for Skill-Based Job Matching, designed to improve fairness and transparency in digital freelancing and recruitment platforms. Traditional ranking systems often depend heavily on historical performance indicators such as ratings, completed projects, reviews, and prior platform activity. Although these signals may help identify experienced candidates, they can also create exposure imbalance by repeatedly prioritizing users who already have visibility. This project addresses that issue by developing a ranking framework that considers both semantic skill relevance and equitable candidate exposure [1].

The proposed system uses Sentence-BERT embeddings to represent freelancer profiles and job descriptions in a semantic vector space. This allows the model to understand contextual similarity between skills and requirements beyond simple keyword matching. A multi-factor scoring mechanism is then used to compute baseline relevance by combining semantic similarity, skill coverage, skill rarity, experience score, and skill depth. This ensures that the ranking process is based on multiple meaningful indicators of candidate-job suitability [3].

To improve fairness, the system applies a post-processing re-ranking framework that includes group-wise score normalization, diversity-aware MMR ranking, and UCB-inspired exposure adjustment. These techniques help redistribute visibility across experience groups while preserving relevance as an important ranking factor. The fairness layer does not randomly promote candidates; instead, it adjusts rankings in a controlled and interpretable manner.

The evaluation results show that the fairness-aware model successfully reduces exposure disparity. The overall exposure disparity decreased from 0.4090 in the baseline model to 0.2425 after fairness adjustment, representing a 40.7% improvement in equity. While ranking accuracy metrics such as NDCG@10 and Precision@10 decreased after fairness intervention, this trade-off is expected because the system balances pure relevance with fairer representation. The results demonstrate that fairness can be incorporated into ranking systems through mathematically grounded and transparent post-processing mechanisms [2],[8].

Overall, the project proves that a job-matching system can be made more inclusive without completely compromising recommendation quality. The model provides a practical, modular, and reproducible foundation for responsible AI-driven recruitment and freelancing platforms.

Key Outcomes

Parameter	Outcome
Fairness Improvement	Exposure disparity reduced from 0.4090 to 0.2425, showing 40.7% improvement in equity.
Representation Balance	Visibility became more evenly distributed across fresher, early-career, mid-level, and senior candidates.
Accuracy Trade-Off	NDCG@10 and Precision@10 decreased after fairness adjustment, showing the measurable cost of fairness-aware ranking.
Transparency	All fairness components, including normalization, MMR selection, and exploration adjustment, are interpretable.
Techniques Used	Sentence-BERT embeddings, cosine similarity, multi-factor scoring, group-wise normalization, MMR re-ranking, and UCB-inspired exposure adjustment.
Scalability	The system is modular, offline-executable, reproducible, and suitable for future API-based integration.

Final Remarks

The Fair Ranking Model for Skill-Based Job Matching demonstrates that algorithmic fairness is both achievable and practically implementable in ranking systems. By combining semantic relevance modeling with fairness-aware re-ranking, the project establishes a balanced approach for improving opportunity distribution in digital employment platforms.

This work confirms that fairness and accuracy do not have to be treated as opposing goals. Instead, they can be managed through a controlled trade-off, where ranking systems remain useful while becoming more inclusive and transparent. The proposed model serves as a scalable foundation for future development of ethical, explainable, and responsible AI-based job-matching systems.

Future Scope

The proposed Fair Ranking Model provides a strong and functional foundation for fairness-aware skill-based job matching. While the current system successfully demonstrates semantic matching, multi-factor ranking, a User feedback signals, such as recruiter interactions, candidate responses, and selection outcomes, can also be incorporated to make the ranking process more adaptive over time and fairness-aware re-ranking in a reproducible offline environment, it can be further extended to support larger-scale and production-oriented use cases.

In future work, the system may be evaluated using real-world datasets from freelancing or recruitment platforms to further validate its effectiveness under practical marketplace conditions.

The framework can be enhanced with context-aware fairness mechanisms, where exposure balancing considers job domain, skill demand, experience level, and role-specific requirements. Knowledge graph-based skill modeling can also be integrated to capture relationships between related technologies and career progression paths, improving the quality of recommendations.

Since the proposed architecture is modular, it can be extended into an API-based or microservice-based deployment for integration with real-time job-matching platforms. These future enhancements would strengthen the system's practical applicability while preserving its core objective of delivering accurate, explainable, and fair ranking outcomes.

The current implementation already establishes the core working pipeline for semantic relevance scoring and fairness-aware exposure balancing. Future extensions would mainly focus on scaling, deployment, and real-world validation rather than changing the fundamental design. This ensures that the proposed model remains a complete final year project while also leaving scope for practical adoption in advanced recruitment and freelancing systems.

REFERENCES

[1] Zou, Jie; Lu, Jie; Zhang, Guangquan; Hu, Fuyuan. *"Fairness and Bias in Recommender Systems: A Survey and Future Directions."* Knowledge-Based Systems, 2023.

[2] Wiesner, Mario; Jugovac, Michael; Jannach, Dietmar. *"A Survey on the Fairness of Recommender Systems."* ACM Computing Surveys, 2022.

[3] Zehlike, Meike; Ghafouri, Soheil; Friedler, Sorelle A. *"Fairness in Rankings: Part II—Learning to Rank and Recommender Systems."* ACM Computing Surveys, 2023.

[4] Biega, E. S.; Zehlike, S.; Gummadi, K. P. *"Fairness-Aware Ranking in Search and Recommendation Systems with Application to LinkedIn Talent Search."* KDD 2021, Knowledge Discovery and Data Mining.

[5] Zhang, Yi; Wei, Wei; Zhang, Yongfeng. *"Fairness and Diversity in Recommender Systems: A Survey."* Springer Nature, 2022.

[6] Mansoury, Mahsa; Nargesian, Fatemeh; Sameki, Mehrnoosh; Waggoner, Bo. *"Joint Multisided Exposure Fairness for Recommendation."* ACM International Conference on Information and Knowledge Management (CIKM), 2022.

[7] Zhou, Zhenyu. *"Fairness among New Items in Cold Start Recommender Systems."* ACM SIGIR Conference on Research and Development in Information Retrieval, 2022.

[8] Ferraro, Nicolas; Lancellotti, Francesco; Martiros, Yannis; Loras, Charles. *"Joint Evaluation of Fairness and Relevance in Recommender Systems with Pareto Frontier."* ACM Recommender Systems (RecSys), 2023.

[9] Wang, Chengyu; Sang, Jitao; Liu, Zhiwei. *"Contrastive Learning for Debiased Gradient Generation in Large-Scale Recommender Systems."* World Wide Web Conference (WWW), 2023.

[10] Mansoury, Mahsa; Valera, Isabel; Ferrara, Antonio; Richardson, Matthew; Bonchi, Francesco. *"A Graph-Based Approach for Mitigating Multi-Sided Exposure Bias in Recommender Systems."* Information Sciences, Elsevier, 2022.

[11] Wu, Wenjie; Yang, Li; Wang, Yunlong. *"Comprehensive Meta-Learned Recommender Systems."* Knowledge-Based Systems, 2023.

[12] Anantharam, Harsha; Rajesh, Shashank; Kumar, Aniket. *"BERT-Based Job Recommendation System Using LinkedIn Dataset."* International Journal of Information Management & Engineering, 2023.

Information Regarding Student

Student Name	Email Id	Permanent Address	Phone Number	Photograph
Trisha K	mahimatrisha2@gmail.com	108, 1st R block, Rajajinagar, Bengaluru North, Bengaluru – 560010, Karnataka	7483070399	
Maitraiyee Bhatt	maitraiyeebhatt@gmail.com	503, Balleza By Lotus, Andheri W, Mumbai-400053, Maharashtra	9372219762	